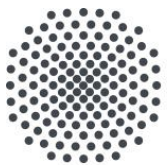


# Introduction to Linux



University of Stuttgart

**Brian Setz, M. Sc.**

[brian.setz@iaas.uni-stuttgart.de](mailto:brian.setz@iaas.uni-stuttgart.de)

Institute of Architecture of Application Systems

# Contents

---

- Introduction
- Brief History
- Overview of Linux
- Debian
- Environment Setup
- Exercises

# Introduction - Lab Sessions

---

- Lab Sessions
  - Focus on Linux, understand how Linux functions
  - Interactive examples and demonstrations
  - Assignments at the end of every lab
- Grading: pass / resubmit / fail
  - Max. 2 fails to pass the lab → requirement for exam
  - Not handing in on time → fail
  - Not handing in as PDF → fail
- Material from labs can (will) be covered on the exam

# Format

- 7 exercises (Lab sessions on Friday at 14:00-15:30 **according to lab schedule slide**) → Exercises in **groups of two**
  - Requires (root) access to Linux (supported: GNU/Linux Debian)
  - No laptop / PC? Find a partner (possibly use AWS Free Tier)
  - Apply theory from lectures in a Linux environment
  - Copying exercises is forbidden, your responsibility to keep your answers private
    - Cite your sources
- **Deadline: see the deadline slide for the exact time and date**
  - Submission: ILIAS → Operating Systems (neue PO) → 0 Exercises → Submission
  - Write a short, coherent report, include lab number, names + student numbers, submit as PDF before the deadline
  - Not only results, include all the steps you took to get to the results
- **Before Monday 28.10.2019 23:59:59, send email to Brian with the two names + student numbers of your group members → [brian.setz@iaas.uni-stuttgart.de](mailto:brian.setz@iaas.uni-stuttgart.de)**

# Introduction – Preliminary Schedule Labs @ 14:00-15:30

Date Lab	Lab Topic
18.10.2019 @ 14:00	<i>No lab</i>
<b>25.10.2019 @ 14:00</b>	<b>1: Introduction, Linux, Virtualization</b>
<b>08.11.2019 @ 14:00</b>	<b>2: Process Management</b>
15.11.2019 @ 14:00	<i>No Lab</i>
22.11.2019 @ 14:00	<i>No Lab</i>
<b>29.11.2019 @ 14:00</b>	<b>3: Process Scheduling</b>
06.12.2019 @ 14:00	Q&A
<b>13.12.2019 @ 14:00</b>	<b>4: System Calls + Interrupts</b>
20.12.2019 @ 14:00	Q&A
<b>10.01.2020 @ 14:00</b>	<b>5: Synchronization + Time</b>
<b>17.01.2020 @ 14:00</b>	<b>6: Memory Management</b>
24.01.2020 @ 14:00	Q&A
<b>31.01.2020 @ 14:00</b>	<b>7: Virtual File System &amp; Block I/ O &amp; I/O Systems</b>
07.02.2020 @ 14:00	Q&A

# Introduction – Preliminary Schedule Deadlines

---

Date Lab	Lab Topic
07.11.2019 @ 23:55:00	Deadline Lab 1
21.11.2019 @ 23:55:00	Deadline Lab 2
12.12.2019 @ 23:55:00	Deadline Lab 3
09.01.2019 @ 23:55:00	Deadline Lab 4
23.01.2019 @ 23:55:00	Deadline Lab 5
30.01.2019 @ 23:55:00	Deadline Lab 6
13.02.2019 @ 23:55:00	Deadline Lab 7

Linux

# What is Linux?

- Linux, free and open-source operating system built around the Linux kernel
- Linux kernel found in:
  - Linux OS's
  - Android
  - Google's Chrome OS
- Pre-emptive multitasking operating system
- Supports symmetrical multiprocessing
- Asynchronous interrupts
- Portable (24+ different processor architectures)





# Linux – Brief History

- 1970 – **Ken Thompson, Dennis Ritchie** (AT&T Bell Lab) release Unix
  - Portable, Modular → Shell scripting + command line
- 1977 – Development of Berkeley Software Distribution (BSD)
- 1983 – **Richard Stallman** starts GNU (GNU is Not Unix) project
  - Goal: free UNIX-like operating system
- 1985 – Intel releases 80386 (i386)
  - First x86 microprocessor with 32 bit instructions + memory management with paging
- 1987 – **Andrew S. Tanenbaum** releases MINIX → academic, sourcecode restricted
- 1992 – BSD vs. Unix lawsuit → limits adaption and development
- 1990-1992 → GNU kernel (GNU Hurd) fails to attract development effort
- 1990's → Lack of a free and widely adopted kernel



**Richard Stallman**

# Linux Kernel

- 1991 – **Linus Torvalds** develops first version of Linux
  - Usenet posting in comp.os.minix

*Hello everybody out there using minix -*

*I'm doing a (free) operating system (**just a hobby, won't be big and professional** like gnu) [...] Any suggestions are welcome, but I won't promise I'll implement them :-)*

*Linus (torvalds@kruuna.helsinki.fi)*

*PS. Yes - it's free of any minix code, and it has a multi-threaded fs. It is NOT portable (uses 386 task switching etc), and it **probably never will support anything other than AT-harddisks**, as that's all I have :-).*

*— Linus Torvalds*

# Linux vs. GNU/Linux

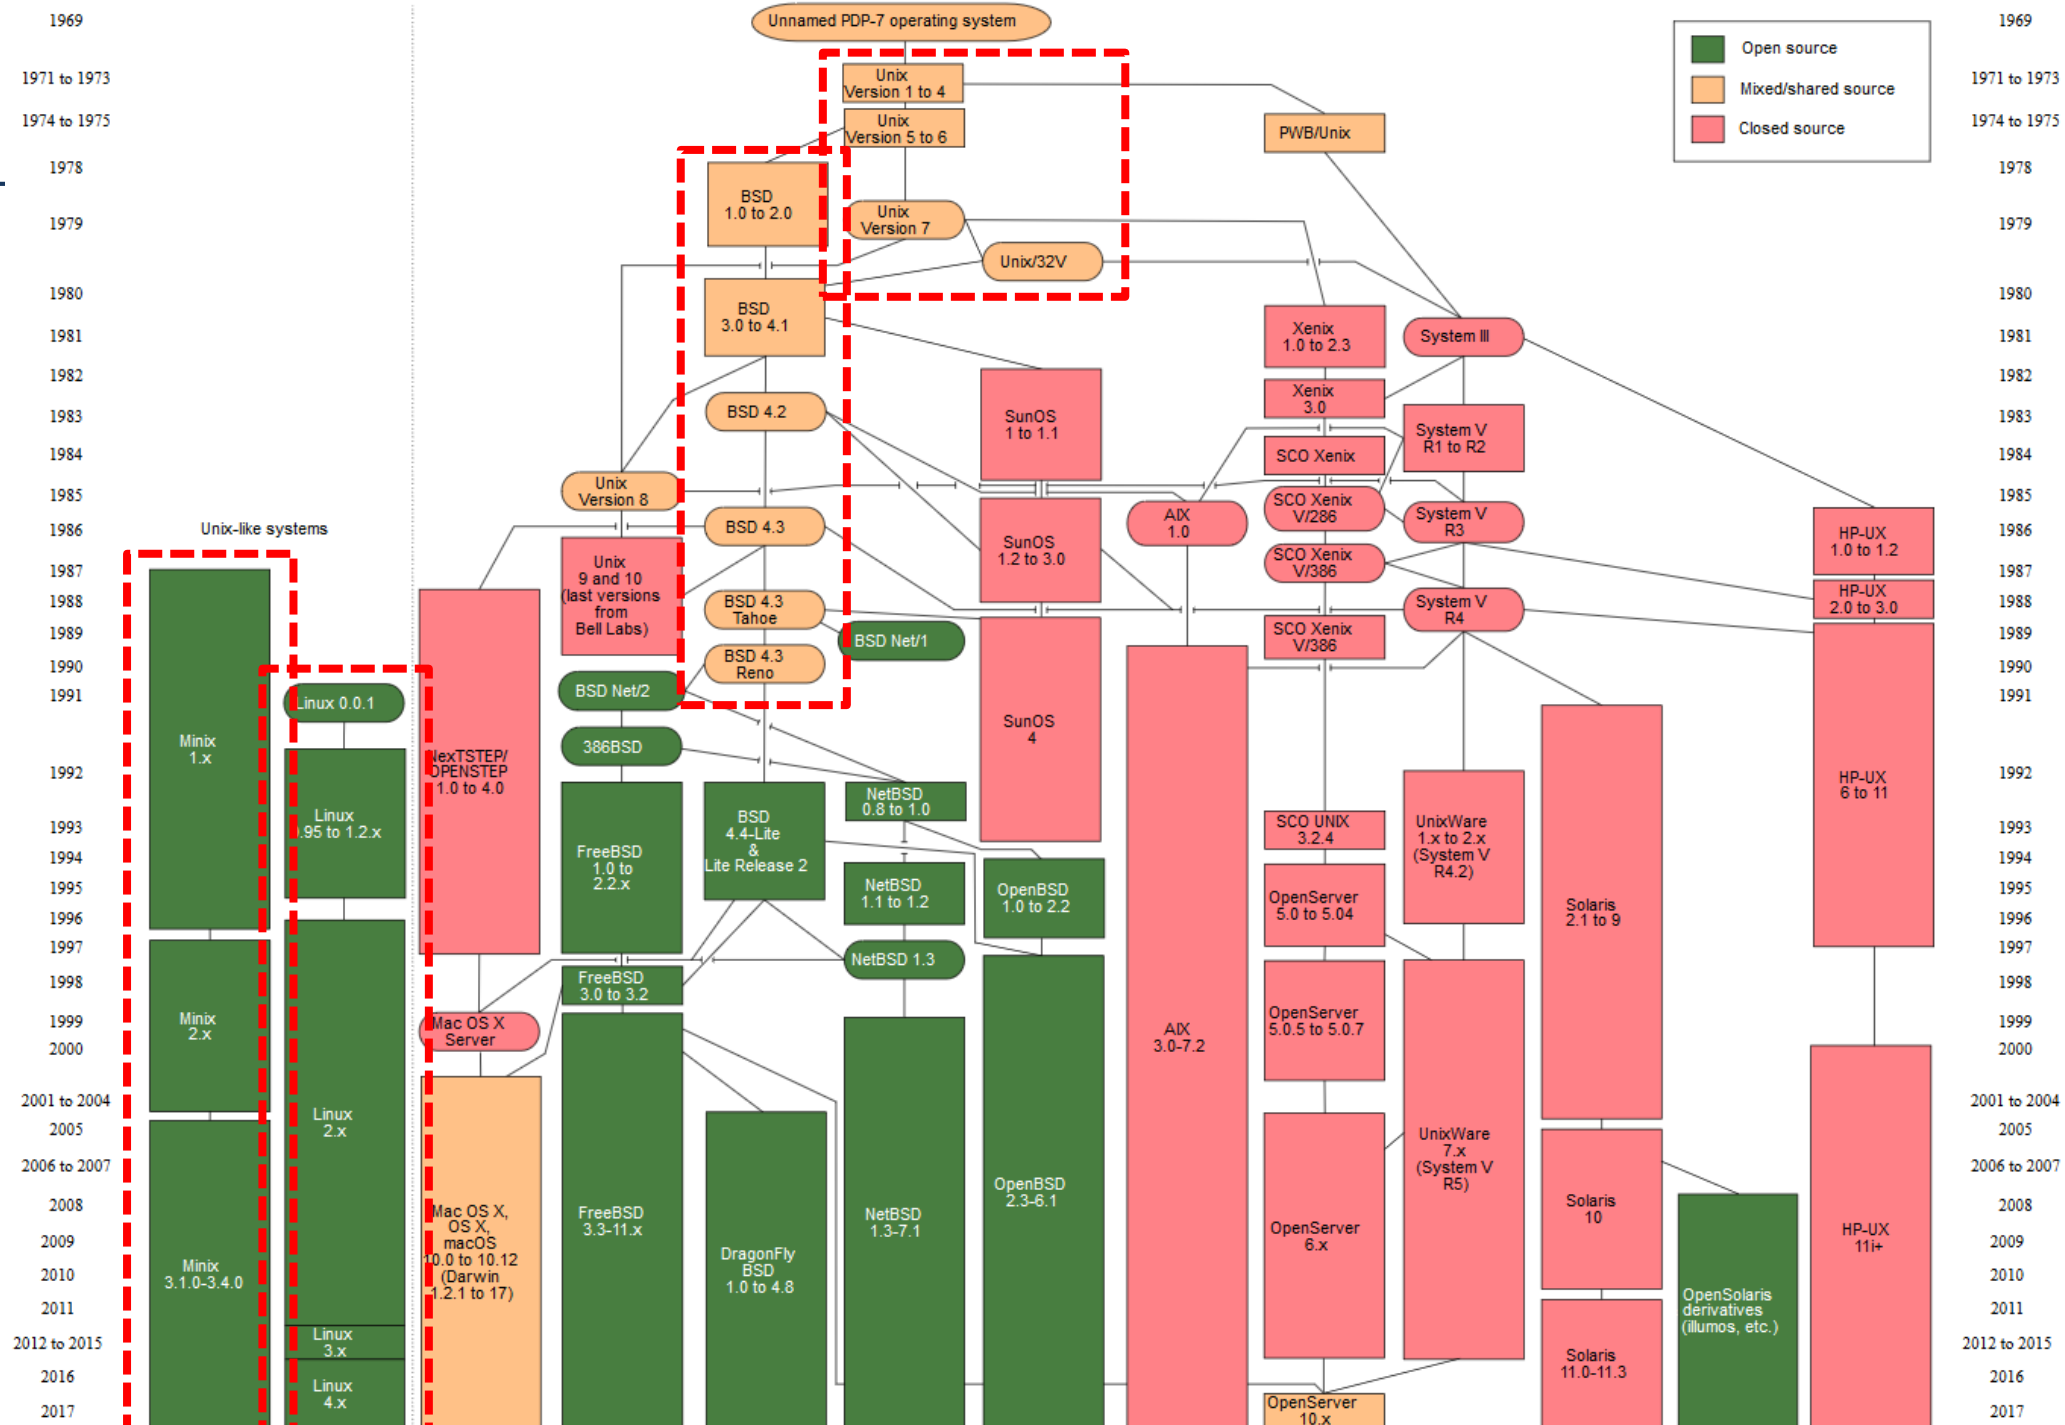
---

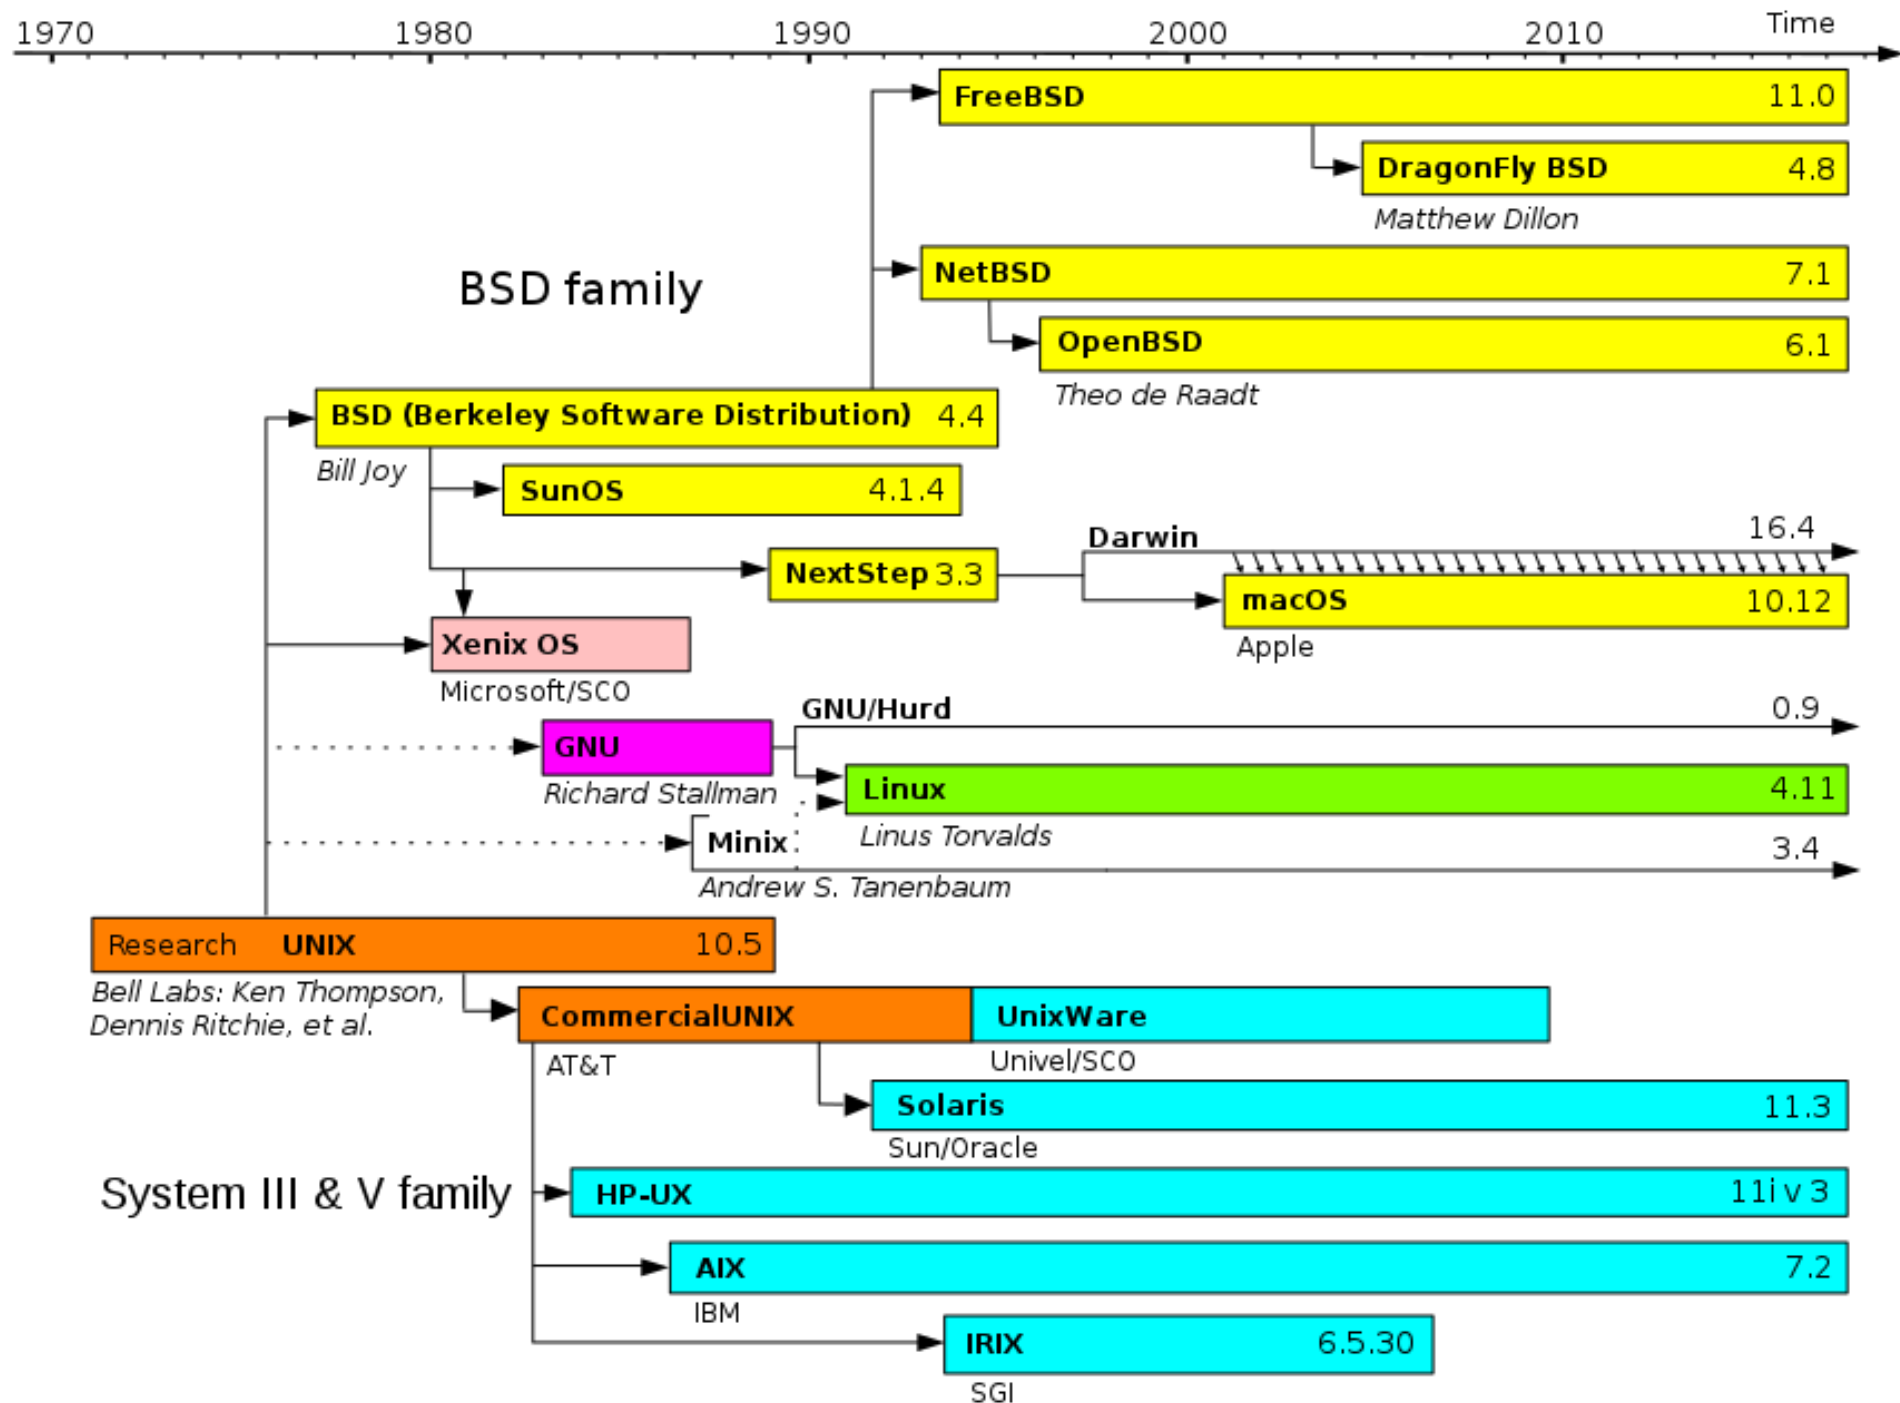
- Linux vs. GNU vs. GNU/Linux?
- Linux → kernel, mediate access to resources such as CPU, RAM, IO (GNU C)
- GNU → tooling, e.g. `bash`, `cat`, `cp`, `mv`, `screen`, `tar`, and many more
- GNU/Linux → Unix-like Operating System
  - Uses Linux kernel and GNU tooling to create a fully functional system
- 1994 – Linux 1.0 → Support for other processor architectures
- 1996 – Linux 2.0 → Symmetric Multiprocessing (SMP)
- 2011 – Linux 3.0 → Drivers and hardware support, 20<sup>th</sup> anniversary
- 2015 – Linux 4.0 → Live kernel patching, file system improvements
- 2019 – Linux 5.0 → Spectre and Meltdown security patches, hardware support
- 15 Sept 2019 – Linux 5.3 → Latest Stable Kernel

# OS Standardization

---

- POSIX → Portable Operating System Interface
  - Initiative of Richard Stallman
  - Standard Application Programming Interfaces (API's) for Unix-like systems
  - Enables creation of tools, applications and platforms that work on a range of OS's
    - Portable code
  - Compliant → macOS, Solaris ...
    - Mostly compliant → Android, FreeBSD, OpenBSD, Linux ...
- LSB → Linux Standard Base
  - Extends POSIX, specifically for Linux distributions
  - Standard file system layout
  - Standard common packages (applications)





# Tanenbaum-Torvalds debate (MINIX vs Linux)

---

**Andrew S. Tanenbaum**



**Linus Torvalds**





# Tanenbaum-Torvalds debate (MINIX vs Linux)

---

- Torvalds is inspired by Tanenbaum's MINIX
  - Tanenbaum's book, Operating Systems: Design and Implementation
- In 1992, Tanenbaum begins a Usenet discussion in `comp.os.minix`
  - Argues micro kernels (MINIX) are superior to monolithic kernels (Linux)
  - Poor portability, too closely tied to Intel's (80)386
  - Linux not suitable for students
- Torvalds replies
  - MINIX has design flaws (no multithreading)
  - Linux is free and developed in spare time
- Controversy about Linux copying MINIX source code
  - Proven false



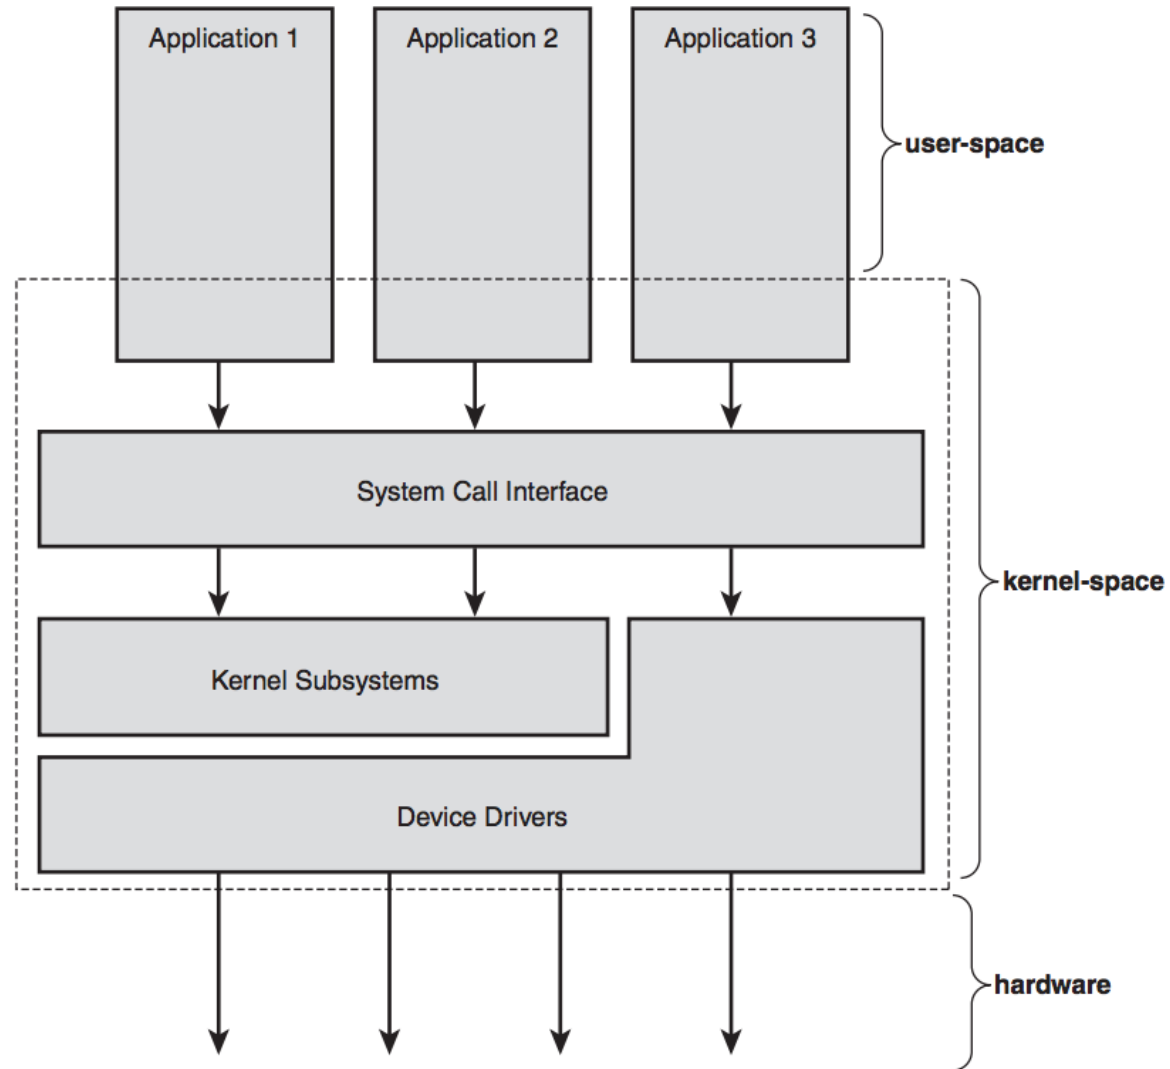
- Kernel includes
  - Interrupt handler → service interrupt requests
  - Process scheduler → share processor time among processes
  - Memory manager → process address spaces
  - System services → networking, inter-process communication
- Kernel executed in elevated system state called **kernel-space**
  - Protected memory space
  - Full access to hardware
  - Kernel threads
- User applications executed in **user-space**

# User-space Kernel-space

---

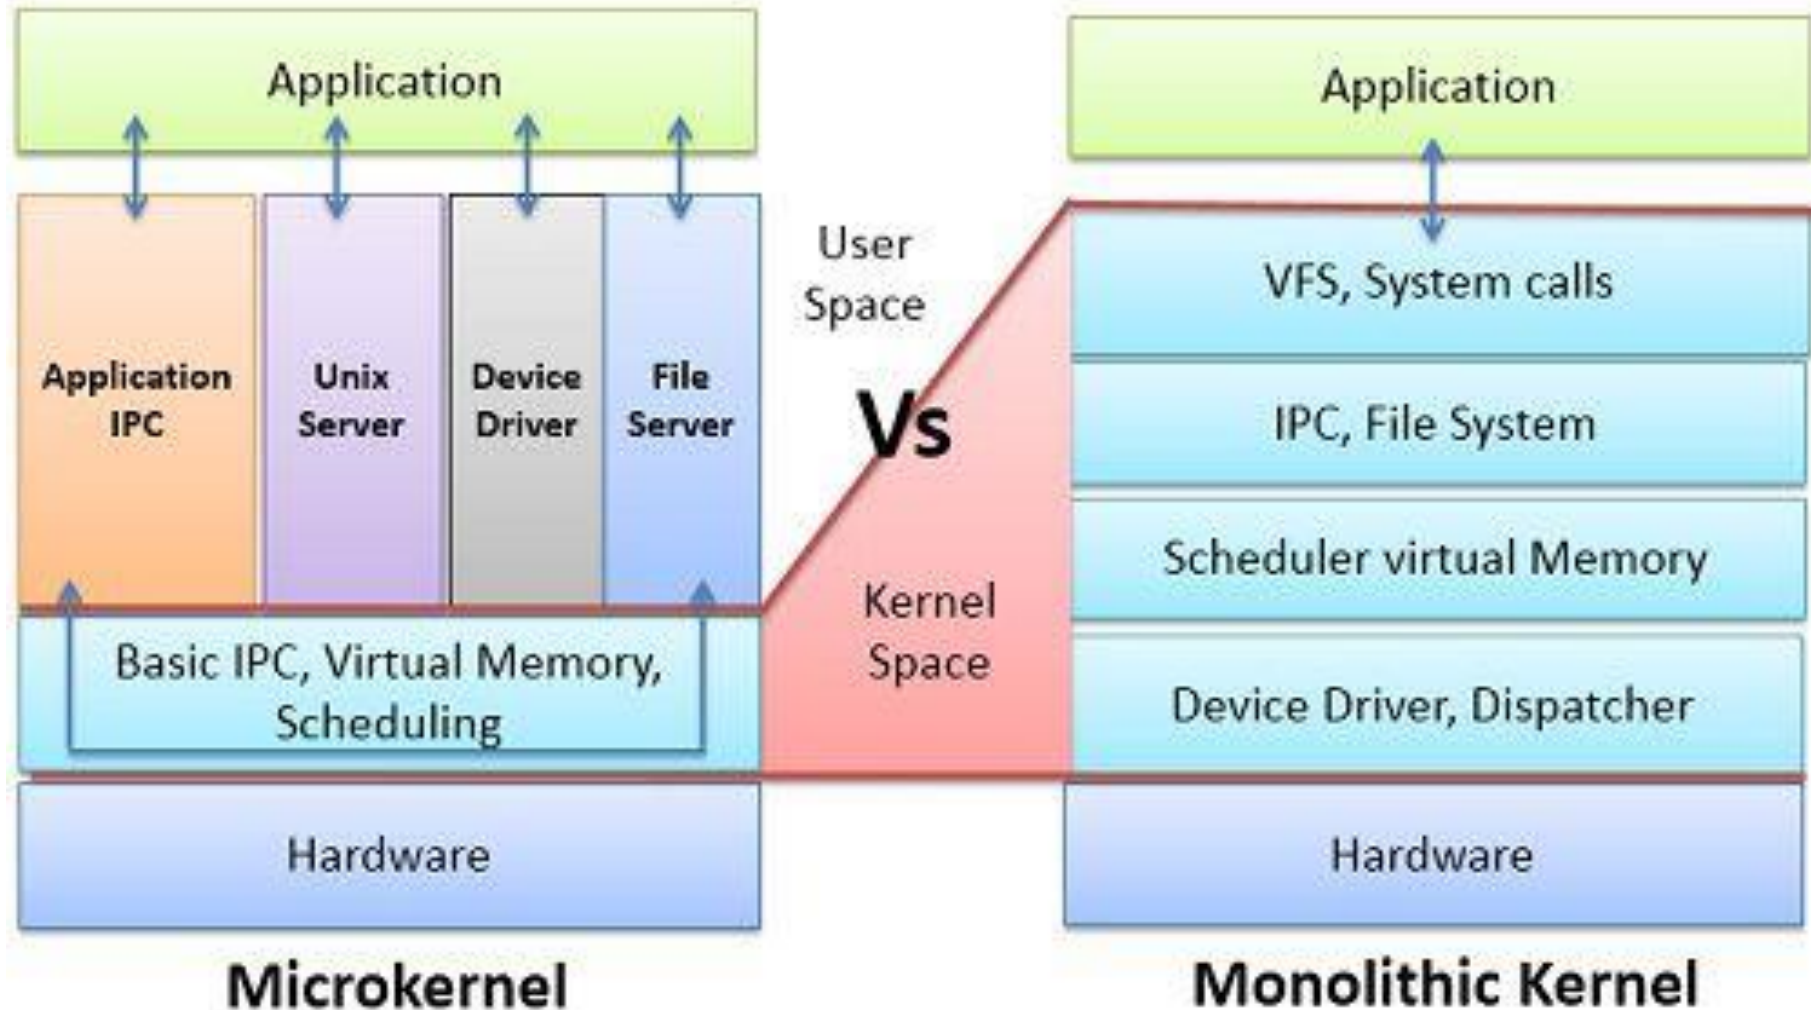
- Applications communicate with kernel via system calls
  - E.g. `accept()`, `chmod()`, `exit()`, `fork()`, `ioctl()`, `recv()`, `write()`
- Kernel communicates with hardware via interrupts
  - Hardware wants to interact → sends interrupt to processor → caught by kernel
  - For synchronization kernel can disable interrupts (all/one)
  - Run in interrupt context, not process context
- In Linux kernels, a processor does one of the following:
  - Execute user code in a process (user-space)
  - Execute system calls on behalf of a process (kernel-space)
  - Handle interrupts not associated with a process (kernel-space)

# User Space vs. Kernel Space



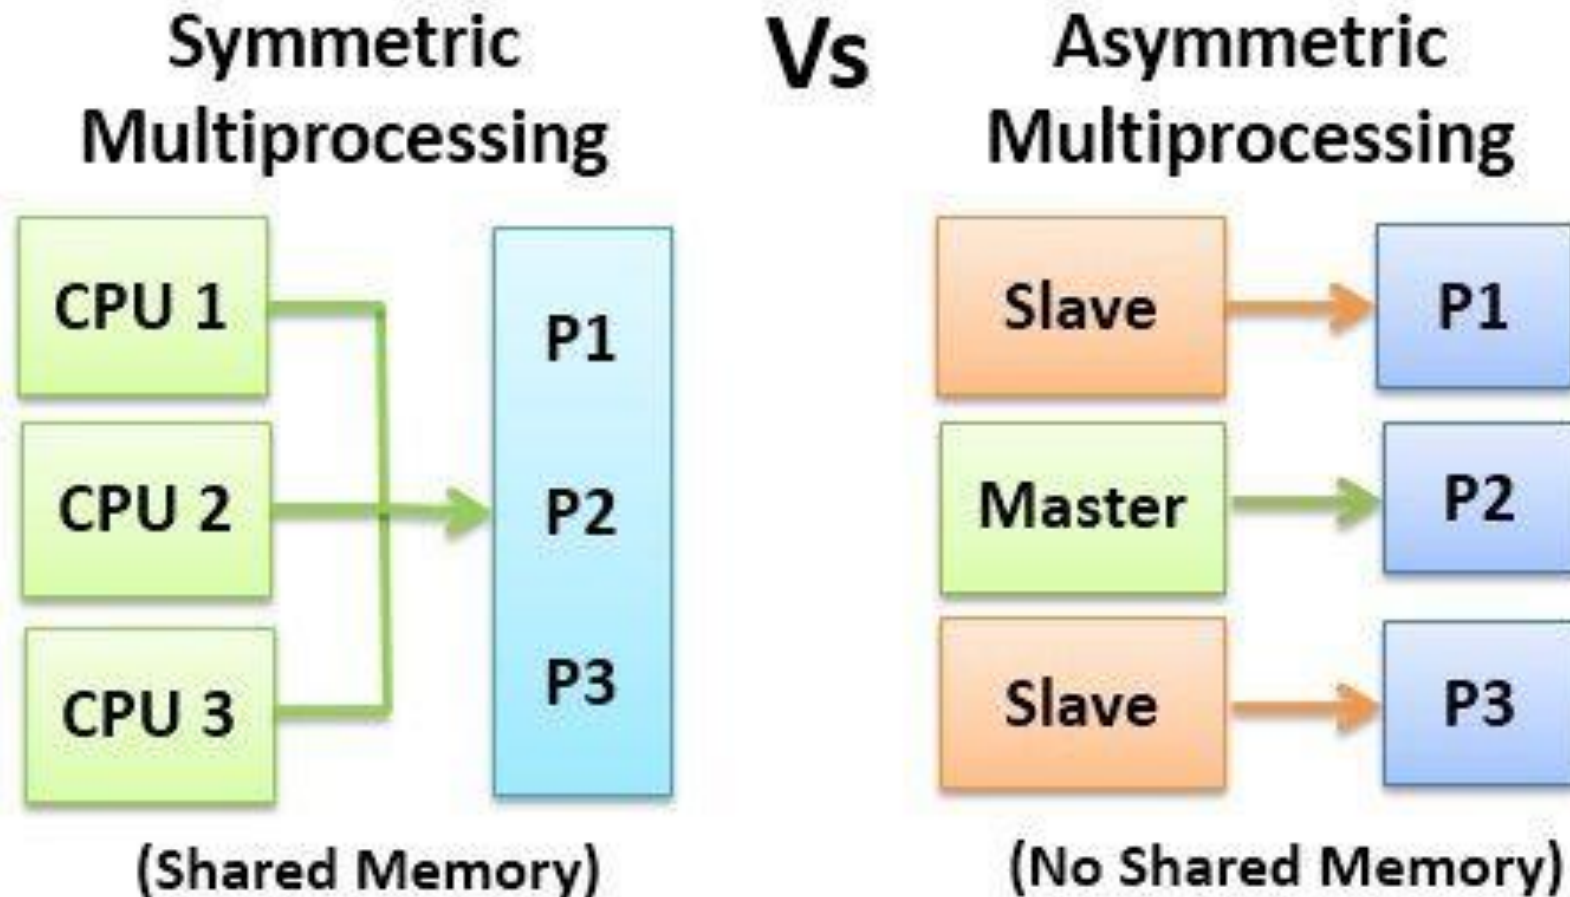
# Linux Kernel characteristics

- Monolithic (as opposed to micro-kernel) → module support



# Linux Kernel characteristics

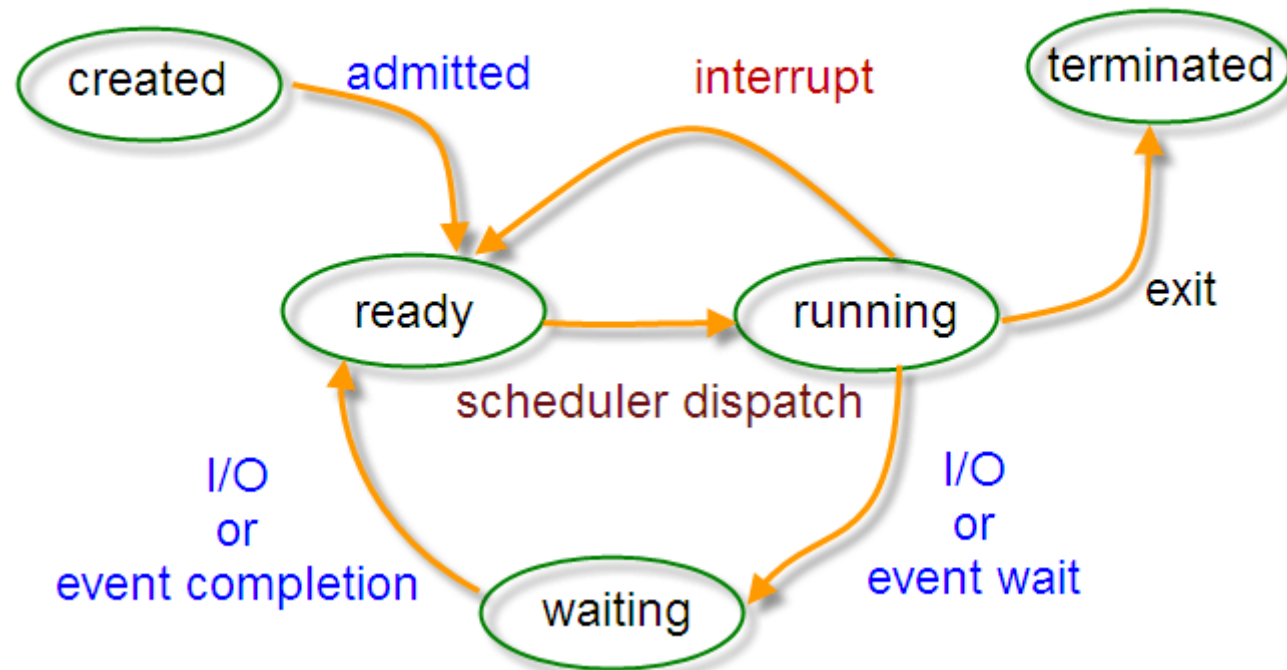
- Monolithic (as opposed to micro-kernel) → module support
- Symmetric Multiprocessor (SMP) support



# Linux Kernel characteristics

- Monolithic (as opposed to micro-kernel) → module support
- Symmetric Multiprocessor (SMP) support
- No differentiation between threads and normal processes, both are tasks

## Process State



# Filesystem Hierarchy Standard (FHS)

---

- `/` → root of the file system
- `/bin/` → Essential binaries (required to boot / rescue) for all users (e.g. bash)
- `/boot/` → boot loader, kernels, initrd
- `/dev/` → devices files
- `/etc/` → system-wide configuration files
- `/home/` → user home directories
- `/lib/` → libraries for binaries in `/bin/` and `/sbin/`
- `/mnt/` → mounted file systems
- `/media/` → removable media
- `/opt/{bin,include,lib,sbin,share}` → application software packages (non-packaged)
- `/proc/` → kernel and process status files

# Filesystem Hierarchy Standard (FHS)

---

- /tmp/ → temporary files
- /sbin/ → Essential system binaries for system administrators (e.g. iptables)
- /sys/ → exporting kernel objects
- /usr/{bin,include,lib,sbin,share,src,X11R6,local,local/bin,local/src} →
  - shareable read-only data (not for boot / rescue), e.g. `man` files
  - `_UNIX` `_system` `_resources`
- /var/ → variable data (logs: /var/log)
- **Non-official**
  - /lost+found/ → recovered files / fragments
  - /selinux/ → SE-Linux settings
  - /srv/ → data served by the system
- [https://refspecs.linuxfoundation.org/FHS\\_3.0/fhs-3.0.pdf](https://refspecs.linuxfoundation.org/FHS_3.0/fhs-3.0.pdf)



# Linux Boot Process

---

- 6 Stages, BIOS-based example

1. BIOS (Basic Input Output System)

- Tests hardware during POST (Power-On Self Test)
- Looks for bootloader in bootable media, loads it into memory, and passes control

2. MBR (Master Boot Record)

- First sector on bootable media (512 bytes)
- Loads and executes GRUB

3. GRUB (Grand Unified Bootloader)

- Read config: `/boot/grub/grub.conf` (BIOS), UEFI slightly different
- Loads and executes kernel (`vmlinux/vmlinuz`) and `initrd` images

# Linux Boot Process

---

## 4. Kernel

- Kernel mounts the temporary root file system from initrd image
- Initrd is used until real root file system is mounted, also contains drivers
  - Newer versions of Linux: initramfs
- Configure RAM and hardware, start process `/sbin/init`

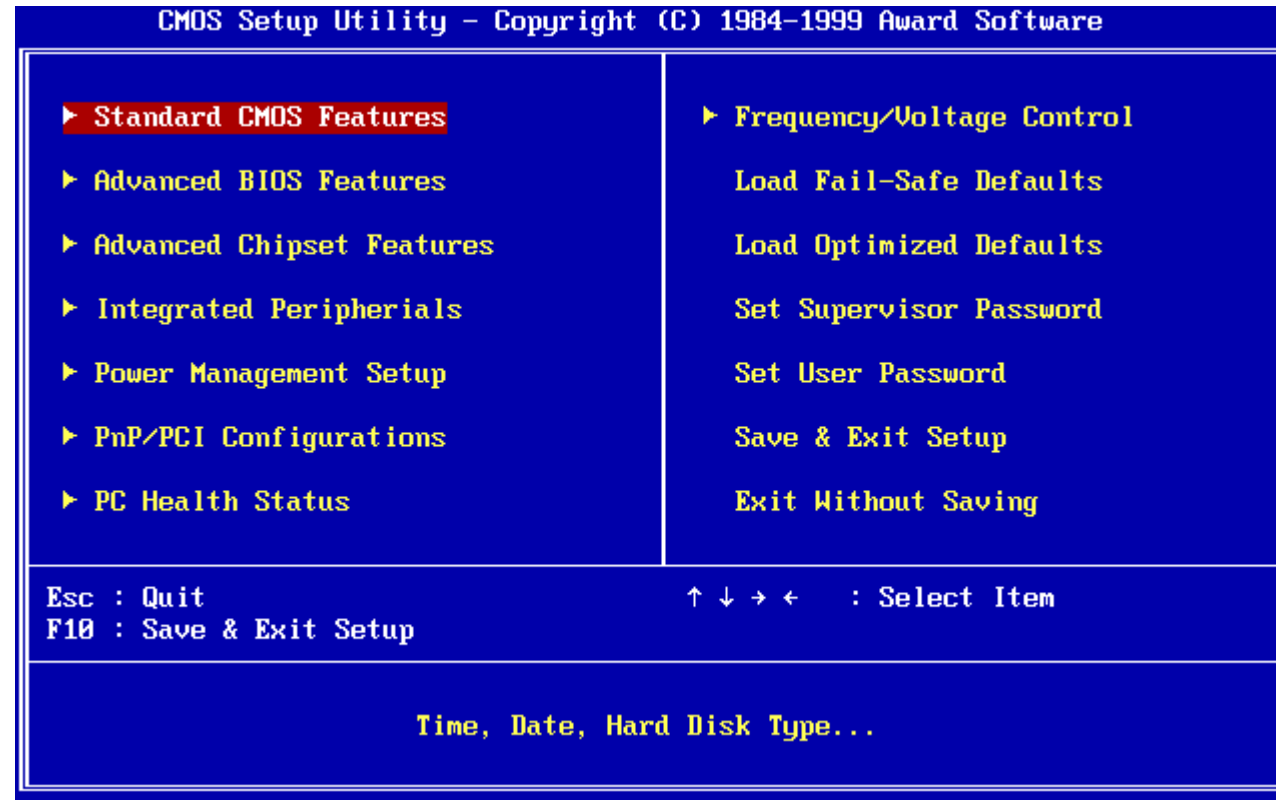
## 5. Init process

- Locate and mount root file system
- Run initialisation scripts (located in `/etc/init.d`), based on runlevel
- Debian: First, run level S → run level 2

## 6. Runlevel applications

- Start services and applications depending on run level
- Levels in Debian
  - 0 - Halt
  - S - Single-user on boot
  - 1 - Single-user
  - 2,3,4,5 - Multi-user
  - 6 – Reboot
  - 7,8,9 - Unused
- Scripts in `/etc/rc<run level>.d/` are executed
  - Symlinks to `/etc/init.d/`

# BIOS



# BIOS vs. UEFI

---

- UEFI
  - Unified Extensible Firmware Interface
  - Faster (fast boot) → caches files for faster start up
  - Allows larger hard drives
    - MBR vs. GUID Partition Table (GPT), MBR supports up to 2.1 TB
  - Processor and driver independent → ARM
  - 32bit and 64bit mode → addressable space
  - Secure boot → malware
- Once booted → GUI Login or Login Shell



# Terminals, Emulators, and Shells

---

- Terminal
- Virtual Terminal
- Pseudo Terminal
- Shell

# Terminal


- Text input/output environment
- Teletypewriter (TTY)
- Hardware-based
- 1960-1980





# Virtual Terminal


- Emulates a hardware terminal
  - CTRL+ALT+F3
- Device files `/dev/tty`



```
Debian GNU/Linux 9 debian tty3
debian login: _
```

# Pseudo Terminal

- Terminal Emulator
  - GNOME Terminal
  - screen
  - ssh
- Device files /dev/pts/



# Shell

---

- Command-line interpreter
- Default login shell: `bash`
- Alternatives, different features
  - `sh`
  - `zsh`
  - `csh`
  - `ksh`
  - ...

- Bourne Shell (sh)
  - Developed by Stephen Bourne (1977)



# Bash

- Bourne Again Shell (bash)
  - Developed by Brian Fox (1989)

```
briansetz@debian:~$
```

# Shell Scripts

---

- Shell scripts are interpreted by the shell
  - Different shells offer different functionality
- How to know which shell should execute a shell script?
- Shebang
  - `#!/bin/sh` → this script will be interpreted by `sh`
  - `#!/bin/bash` → this script will be interpreted by `bash`

# (bash) Shell Commands

- `man` (manual) command
  - Describes how to use commands → documentation
  - How to use `man`? `man man`
- Pipes and redirects
  - Pipe `x | y` → (standard) output of `x` is used as (standard) input for `y`
    - `ps aux | grep sh`
  - Redirect `x [1/2/&]> [file descriptor]` → redirect output of `x` to file descriptor (1 = std. out, 2 = std. err, & = both)
    - `cat /var/log/syslog > /home/briansetz/outlog`
    - `>` → truncate, `>>` → append
- Execute in background → `x &`
- Multiple commands → `x; y; z` or `x && y && z`

# Linux Distributions

---

- List of all GNU/Linux distributions
  - [https://upload.wikimedia.org/wikipedia/commons/1/1b/Linux\\_Distribution\\_Timeline.svg](https://upload.wikimedia.org/wikipedia/commons/1/1b/Linux_Distribution_Timeline.svg)
- Majority (80-90%) of all distributions are based on / derived from:
  - Debian (1993)
  - Slackware (1993)
  - RedHat (1995)
- Some examples...




# Linux Distributions

- Debian



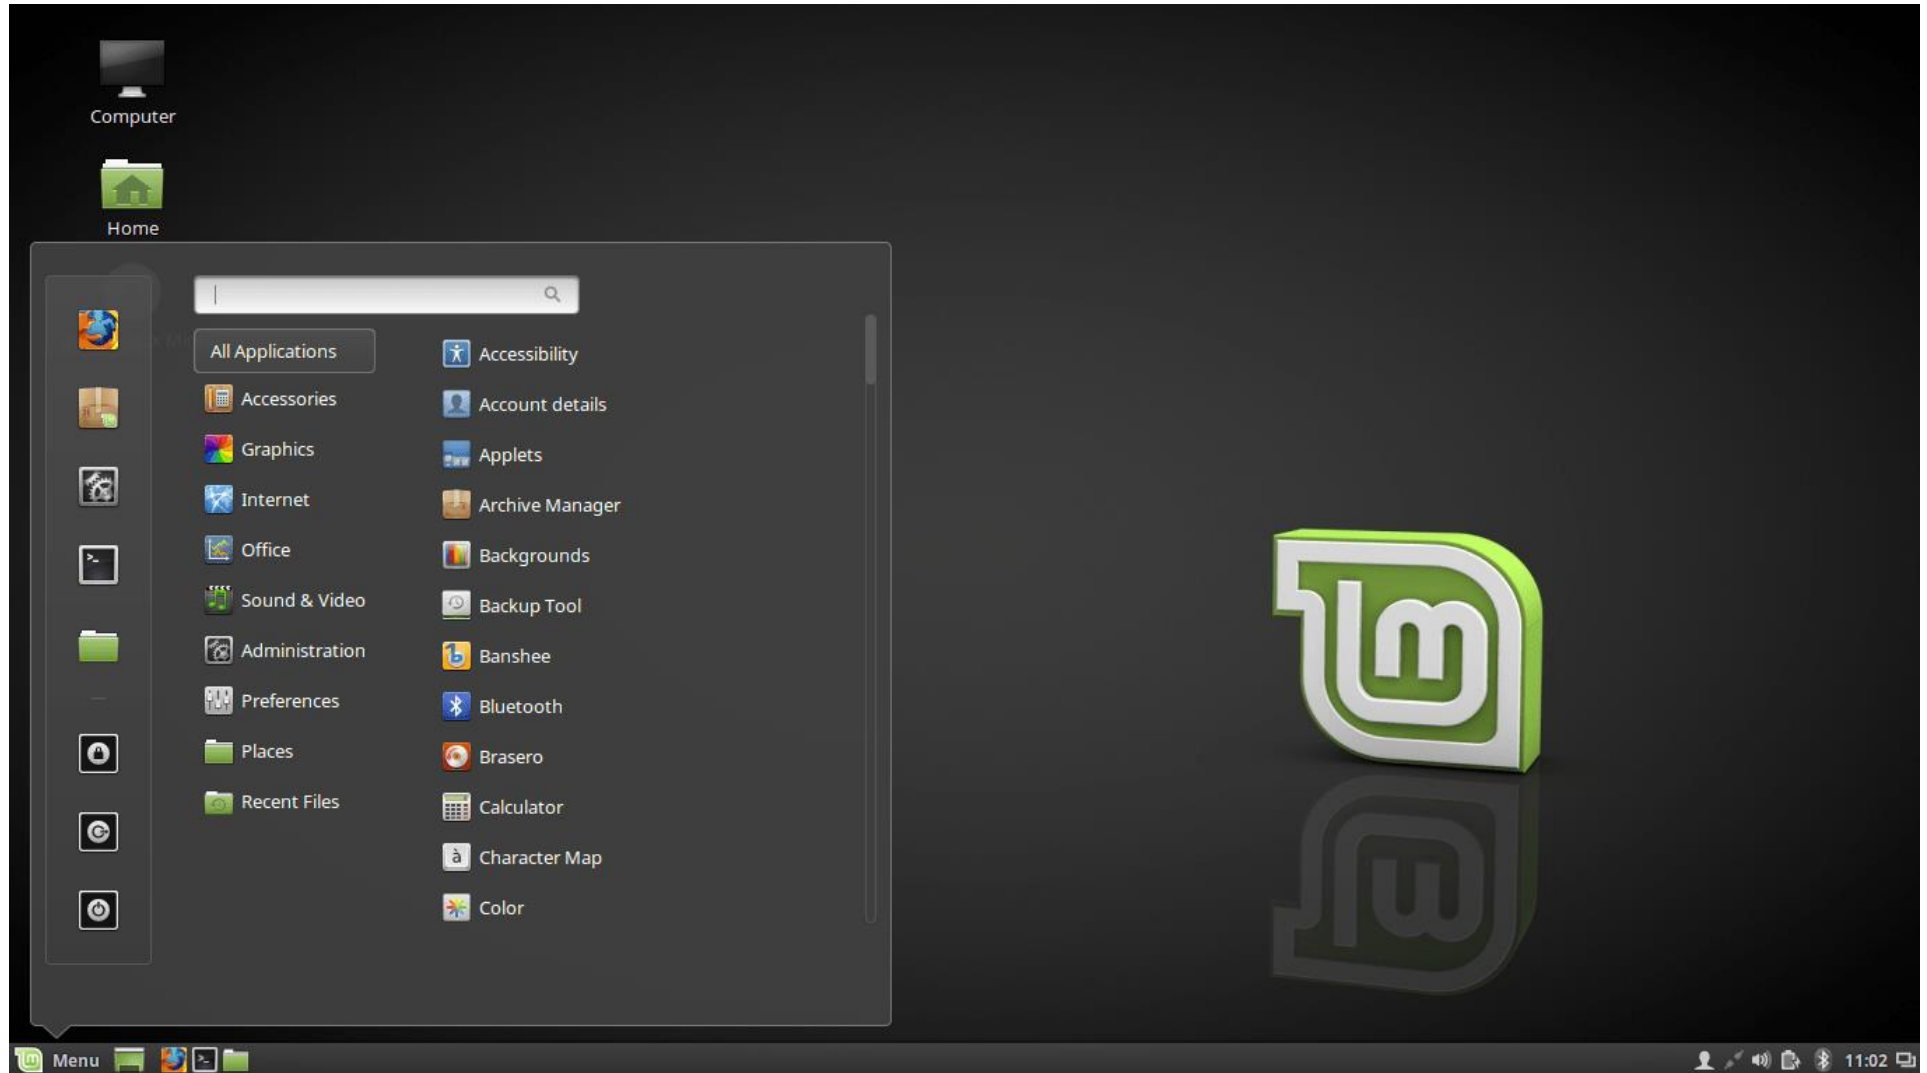
# Linux Distributions

## ■ Ubuntu




# Linux Distributions

- Linux Mint



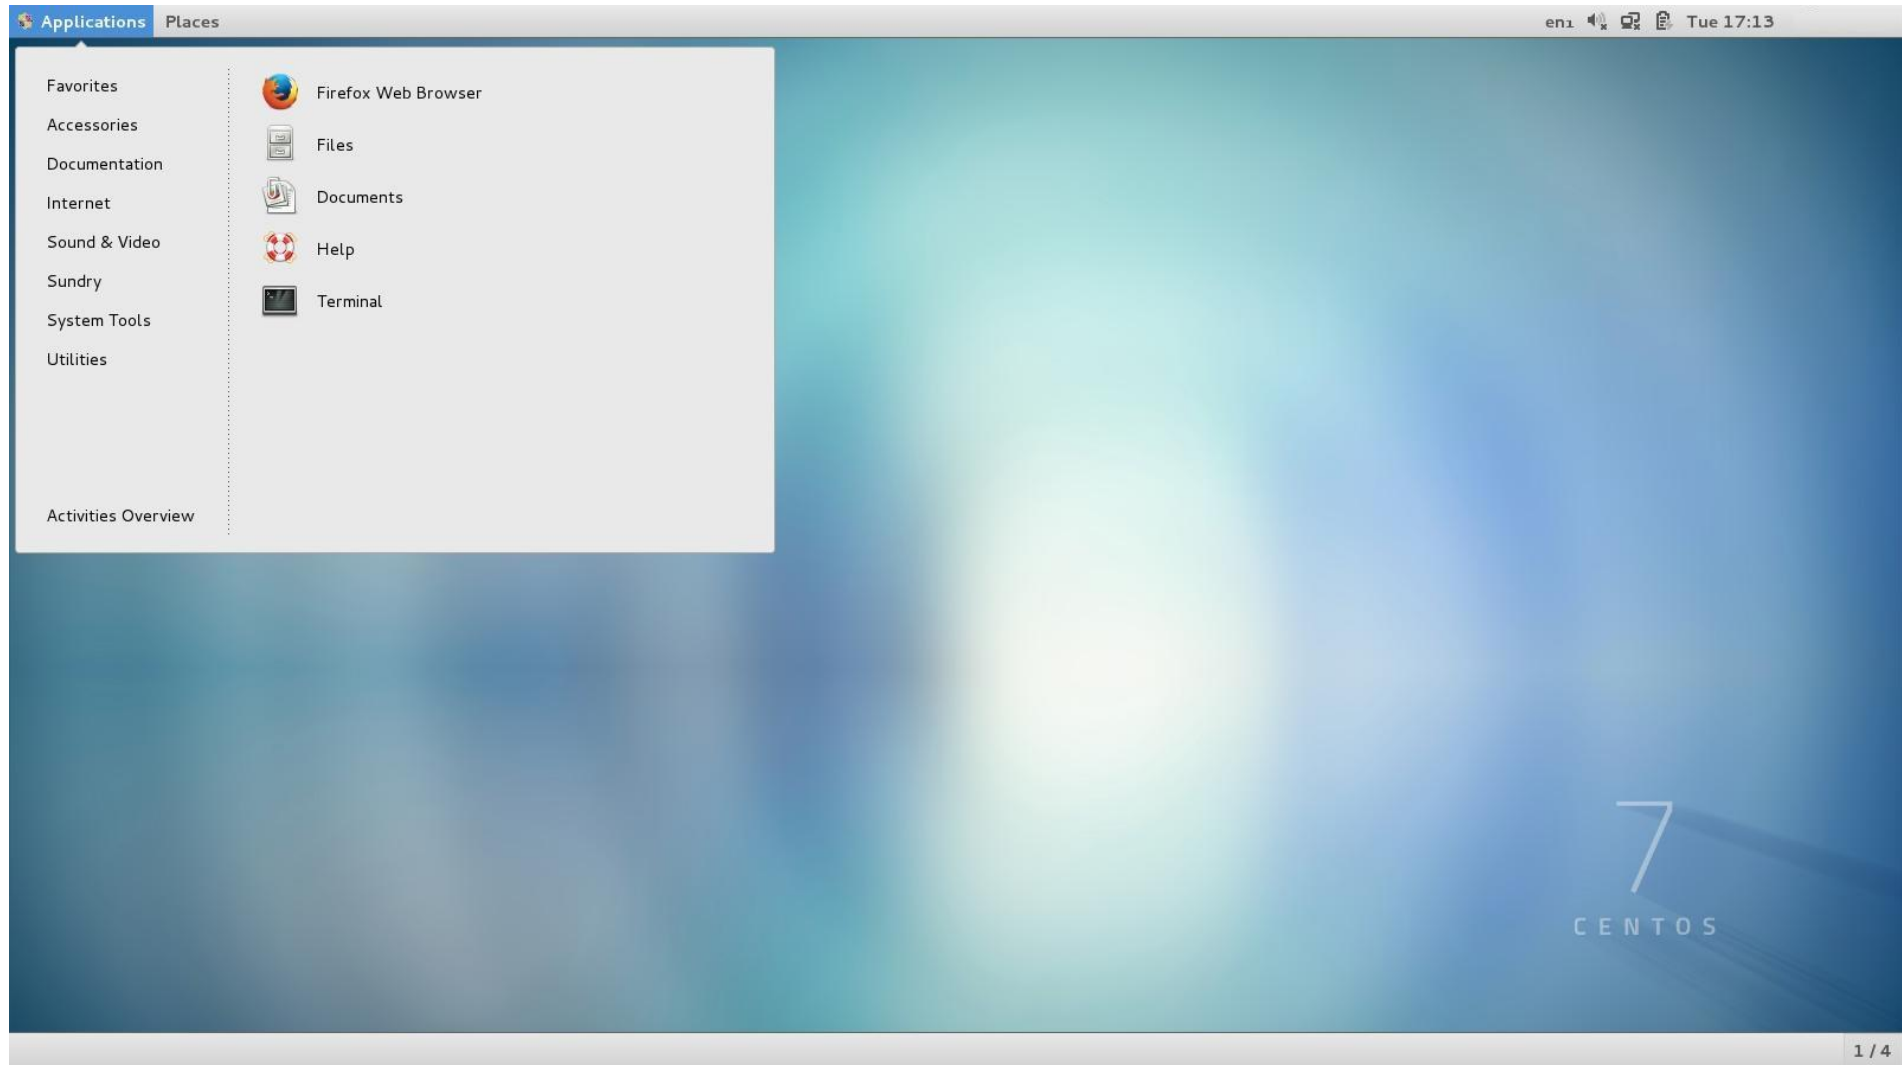
# Linux Distributions

- Elementary OS




# Linux Distributions

## ■ CentOS



# Linux Distributions

- Arch





# Exercises Environment

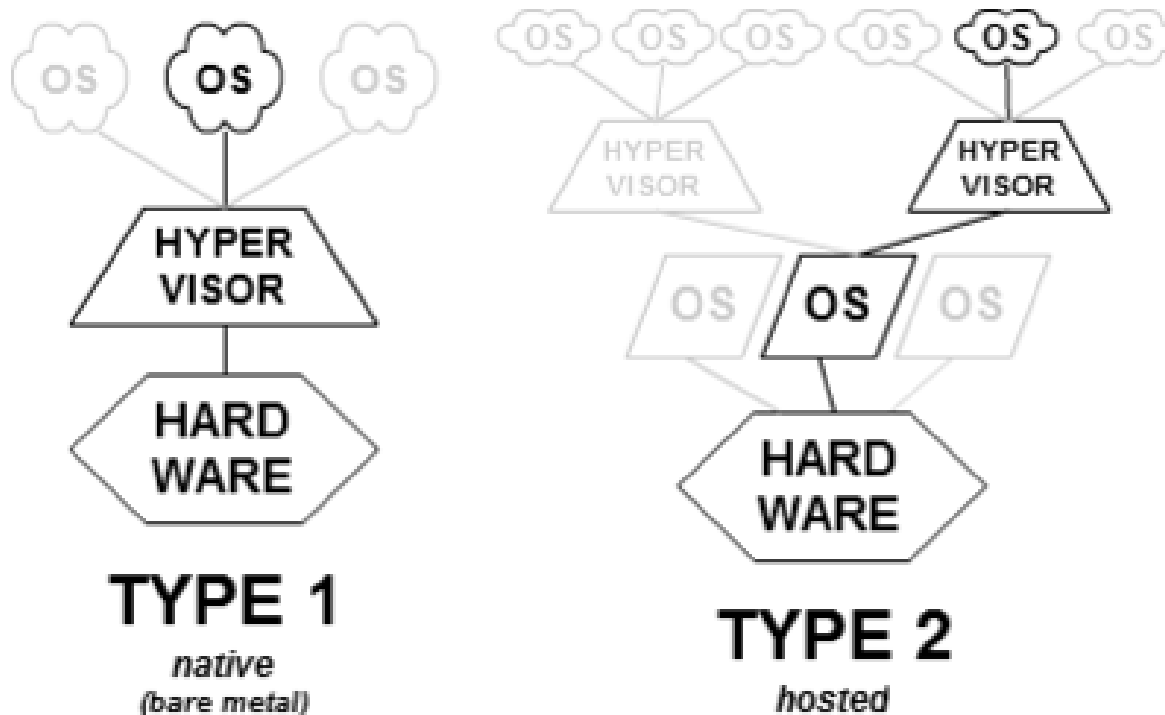


- Unix-like operating system
- Supports multiple kernels: Linux, BSD-based, GNU Hurd
- GNU Tooling (`coreutils`)
- Download: <https://cdimage.debian.org/mirror/cdimage/archive/9.5.0/i386/iso-dvd/>
  - Only download `debian-9.5.0-i386-DVD-1.iso`
  - During labs only Debian 9.5.0 i386 is supported



# Virtualbox

- VirtualBox → Hypervisor, creates and runs virtual machines
  - Type 2 hypervisor
- Download: <https://www.virtualbox.org/wiki/Downloads>
  - Enable VT-d, optionally VT-x also (BIOS)
  - Disable Hyper-V (MS Window's type 1 hypervisor)



# Steps (1)

---

1. Install VirtualBox
2. Create new Virtual Machine
  1. Type: Debian 32bit
  2. Memory: 1GB+
  3. Create a virtual hard disk now → VDI → Dynamically Allocated: 32GB
3. Change Virtual Machine Settings
  1. System → Processor → Increase processors to 50%
  2. Display → Video Memory: Max → Enable 3D Acceleration (optional)
  3. Storage → CD → Choose Virtual Optical Disk File → Select Debian ISO
4. Start the Virtual Machine

## Steps (2)

---

### 5. Install Debian, important settings:

- Use your full name as user name
- Partition disks: **Guided – use entire disk and set up LVM**
- Partitions disks: **Separate /home, /var, and /tmp partitions**
- Configure the package manager: **Scan another CD or DVD? → NO**
- Configure the package manager: **Use a network mirror? → YES**
- Software selection: **GNOME, print server SSH server, standard system utilities**
- Install the GRUB boot loader to the master boot record? → YES

## Steps (3)

---

### 6. Update Debian

- Open Terminal
- Switch to root user: `su`
- Edit `sources.list`: `nano /etc/apt/sources.list` and put a `#`-sign in front of the line starting with `deb cdrom`
- Update the system: `apt update` followed by `apt upgrade`
- Reboot
- Open Terminal, switch to root user, and run: `uname -r` the output should be version `4.9.0-11-686` or greater

### 7. Install dependencies

- As root, run: `m-a prepare` and install the dependencies

## Steps (4)

---

### 8. Install the VirtualBox guest additions

- In VirtualBox, click on Devices → Insert Guest Additions CD image
- In Debian, click cancel and open the File Manager
- Click on the CD, and right click in the window that opened and select Open Terminal
- As root, run: `sh VBoxLinuxAdditions.run`
- Wait for the process to finish

### 9. Configure VirtualBox

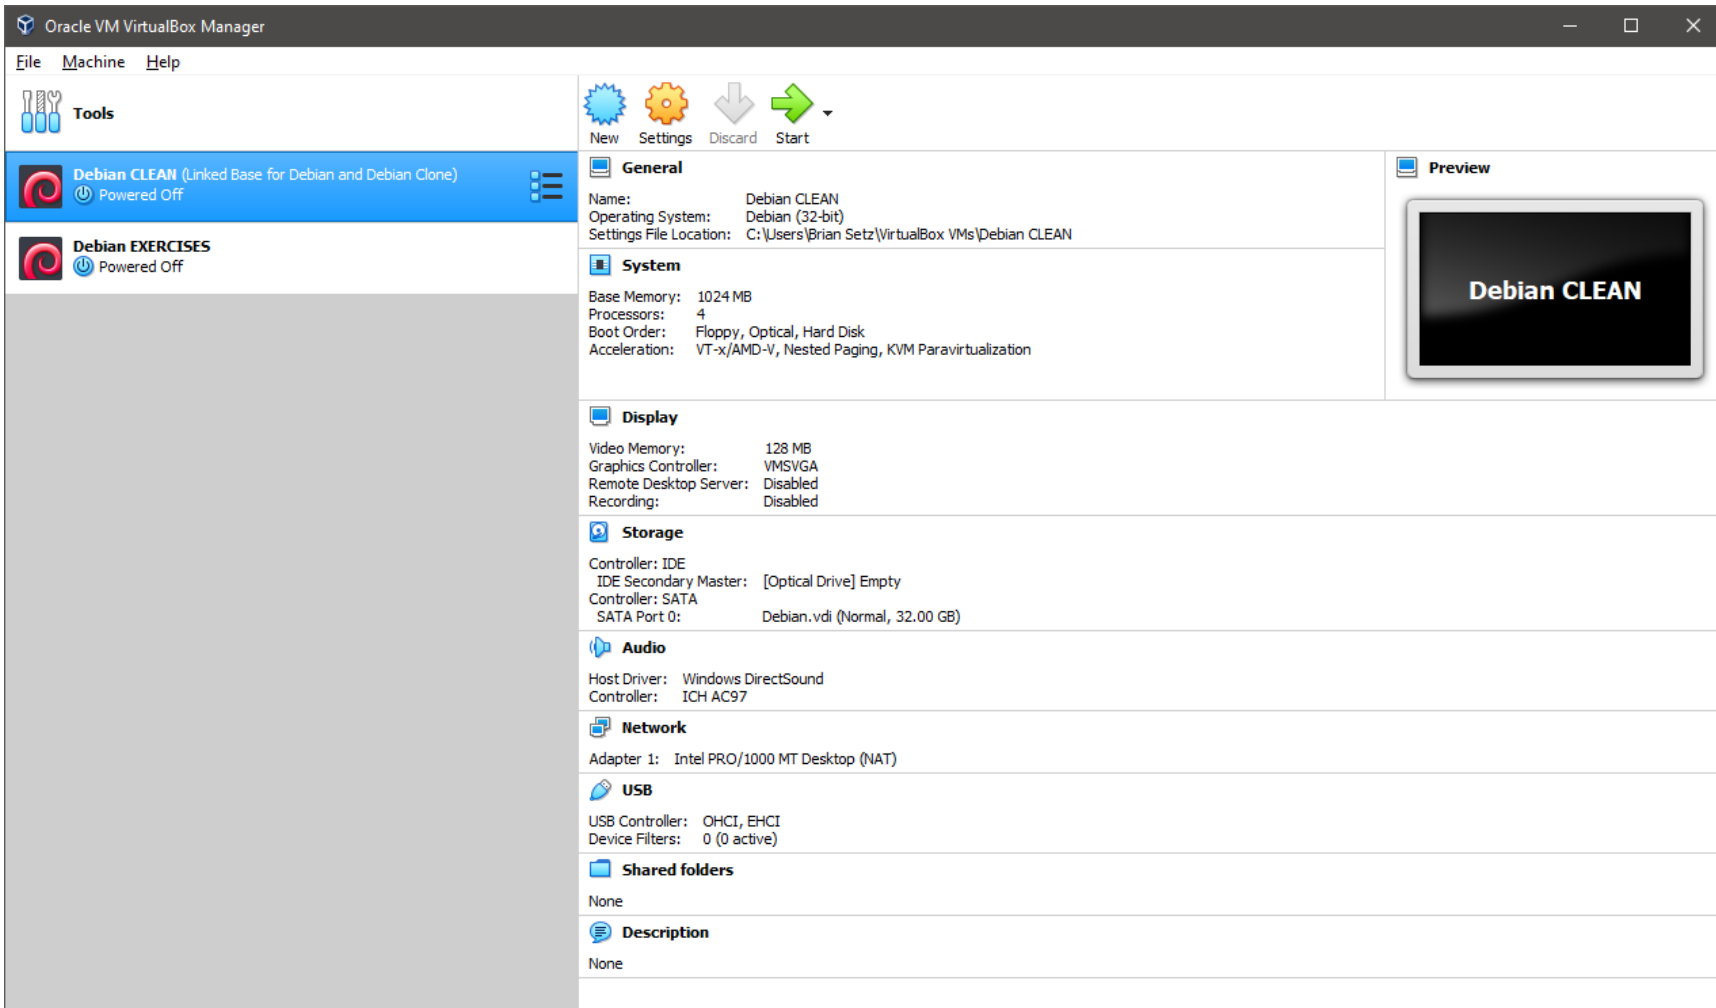
- In VirtualBox click Devices → Shared Clipboard → Bidirectional
- And, Devices → Drag and Drop → Bidirectional
- Shutdown Debian

### 10. In VirtualBox, right click the VM and select Clone → Linked Clone!

- Rename original to: Debian CLEAN, rename clone to: Debian EXERCISES

# Video Instructions

- Video instructions are available on YouTube: [https://youtu.be/S3D\\_2ytNqEs](https://youtu.be/S3D_2ytNqEs)
  - Short version, actual installation takes 30-45 minutes depending on hardware



# Exercises

# Exercises

Always include a full description (e.g. input / output) of the performed tasks, in addition to the answers. Submit a short, coherent report as PDF.

1. Install VirtualBox, setup Debian within VirtualBox, **(a)** use the `lsb_release` command to print all the distribution details, and **(b)** use the `uname` command to print the kernel *release* version. Read the `man` pages of these commands.
2. Determine **(a)** the shell that is used by default by using the `echo` command and the `$SHELL` variable, and **(b)** list the directories found in the `$PATH` variable.
3. List the number of scripts that run **(a)** at run level S, **(b)** run level 2, and **(c)** run level 5. Use the `ls`, `wc` and pipe commands. Hint: check the `man` pages to list each file on a new line and 'pipe' the results from one command to the other.
4. Research the difference between `systemd` and `init`, **(a)** describe in your own words the difference between these systems, and **(b)** determine which of the two is used by the operating system you have installed in VirtualBox. How can you tell?