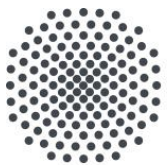


Processes



University of Stuttgart

Brian Setz, MSc

brian.setz@iaas.uni-stuttgart.de

Institute of Architecture of Application Systems

Introduction - Lab Sessions

- Lab Sessions
 - Focus on Linux, understand how Linux functions
 - Interactive examples and demonstrations
 - Assignments at the end of every lab
- Grading: pass / resubmit / fail
 - Max. 2 fails to pass the lab → requirement for exam
 - Not handing in on time → fail
 - Not handing in as PDF → fail
- Material from labs can (will) be covered on the exam

Format

- 7 exercises (Lab sessions on Friday at 14:00-15:30 **according to lab schedule slide**) → Exercises in **groups of two**
 - Requires (root) access to Linux (supported: GNU/Linux Debian)
 - No laptop / PC? Find a partner (possibly use AWS Free Tier)
 - Apply theory from lectures in a Linux environment
 - Copying exercises is forbidden, your responsibility to keep your answers private
 - Cite your sources
- **Deadline: see the deadline slide for the exact time and date**
 - Submission: ILIAS → Operating Systems (neue PO) → 0 Exercises → Submission
 - Write a short, coherent report, include lab number, names + student numbers, submit as PDF before the deadline
 - Not only results, include all the steps you took to get to the results
- **Before Monday 28.10.2019 23:59:59, send email to Brian with the two names + student numbers of your group members → brian.setz@iaas.uni-stuttgart.de**

Introduction – Preliminary Schedule Labs @ 14:00-15:30

Date Lab	Lab Topic
18.10.2019 @ 14:00	<i>No lab</i>
25.10.2019 @ 14:00	1: Introduction, Linux, Virtualization
08.11.2019 @ 14:00	2: Process Management
15.11.2019 @ 14:00	<i>No Lab</i>
22.11.2019 @ 14:00	<i>No Lab</i>
29.11.2019 @ 14:00	3: Process Scheduling
06.12.2019 @ 14:00	Q&A
13.12.2019 @ 14:00	4: System Calls + Interrupts
20.12.2019 @ 14:00	Q&A
10.01.2020 @ 14:00	5: Synchronization + Time
17.01.2020 @ 14:00	6: Memory Management
24.01.2020 @ 14:00	Q&A
31.01.2020 @ 14:00	7: Virtual File System & Block I/ O & I/O Systems
07.02.2020 @ 14:00	Q&A

Introduction – Preliminary Schedule Deadlines

Date Lab	Lab Topic
07.11.2019 @ 23:55:00	Deadline Lab 1
<u>21.11.2019 @ 23:55:00</u>	Deadline Lab 2
12.12.2019 @ 23:55:00	Deadline Lab 3
09.01.2019 @ 23:55:00	Deadline Lab 4
23.01.2019 @ 23:55:00	Deadline Lab 5
30.01.2019 @ 23:55:00	Deadline Lab 6
13.02.2019 @ 23:55:00	Deadline Lab 7

Additional Notes

- Only 1 submission per group
- Grading (pass / fail / resubmit) will be done within 2 weeks
- Read assignments and instructions carefully
 - Always include all steps
 - When implementation is required, include code
 - Not clear? Send an email
- Questions

Agenda

- Previously
 - Unix-Like History, GNU/Linux, Linux Kernel, Terminal
 - Virtual Machine, Virtual Box, Debian
 - Exercises Lab 1
- Today
 - Processes
 - Process descriptor and task structure
 - Process creation
 - Threads vs. Processes
 - Process termination
 - Crash-Course Kernel Modules

Processes

Processes Definitions


- Process → performs tasks in the operating system
- Program → Set of machine code instructions and data
- A process is a program that is executing
- Processes are isolated, have their own address space and cannot interact with other processes except through system calls
- Processes use many different system resources, needs to be tracked and managed
- Kernel has the task to manage processes and their resources

Kernel Process Management

- For each process, the kernel manages at least:
 - Open files
 - File descriptor count
 - Pending signals
 - SIGINT → interrupt process, SIGKILL → kill process, SIGCHLD → Child process stopped
 - Internal kernel data
 - Mappings to kernel addresses
 - Process state
 - Created, running, waiting, exited
 - Memory address space (mappings)
 - Virtual Memory (memory sandbox), page table
 - One or more threads
 - Track children, thread group IDs

Process Identifiers



- In Linux processes are identified by a numeric value called the process ID
- pid → process ID
 - Each process has a unique identifier
 - Defined as an `int` in the `task_struct`, max value: 32 768 (`short int`)
 - `cat /proc/sys/kernel/pid_max`
 - Limits process count → relevant for servers
- ppid → parent process ID: The process identifier of the parent process
- tgid → thread group ID: The ID of the process that started the thread
- tid → thread ID: The ID of the thread, unique for each thread
- Process has one thread? `pid == tgid == tid`
- Process has multiple threads? Same tgid for all threads, `tgid == pid` of first thread, unique tid per thread



init is the grandfather of all processes.

All processes running in Linux can trace their relationship back to init.

Process1 is a child of init and a parent of Process2.






Process3 is both a child and a parent. It is a child process of init and a parent of processes 4 and 5.



A process may have any number of child processes, but it may have only one parent process.

The PPID of a child is the same as the PID of its parent.

Process2 has a PPID of 101 because its parent is Process1, which has a PID of 101.



Process2 is a child of Process1.



Process lineage may extend to multiple levels.

Each PID is unique, but duplicate PPIDs are allowed since a parent process may have several child processes.

Threads

- Thread → objects of activity within a process, execution flow
 - Program counter
 - Stack
 - Registers
- The Linux kernel schedules tasks NOT processes OR threads
 - Traditional UNIX → one process = one thread
 - Modern Linux → one process = at least one thread
- Linux threads are processes (tasks) that **can** share resources (e.g. memory, file descriptors, signals) → described in the process descriptor

Processes in Linux

(process id, thread group id)

- `ps tree` command to view (sub) process tree

```
briansetz@debian:~$ ps tree -pg 1765
gnome-terminal-(1765,1765)─bash(1827,1827)─pstree(2147,2147)
                             │
                             ├──{dconf worker}(1798,1765)
                             ├──{gdbus}(1782,1765)
                             └──{gmain}(1781,1765)
```

Shell

pstree process

Terminal Emulator

GNOME application settings

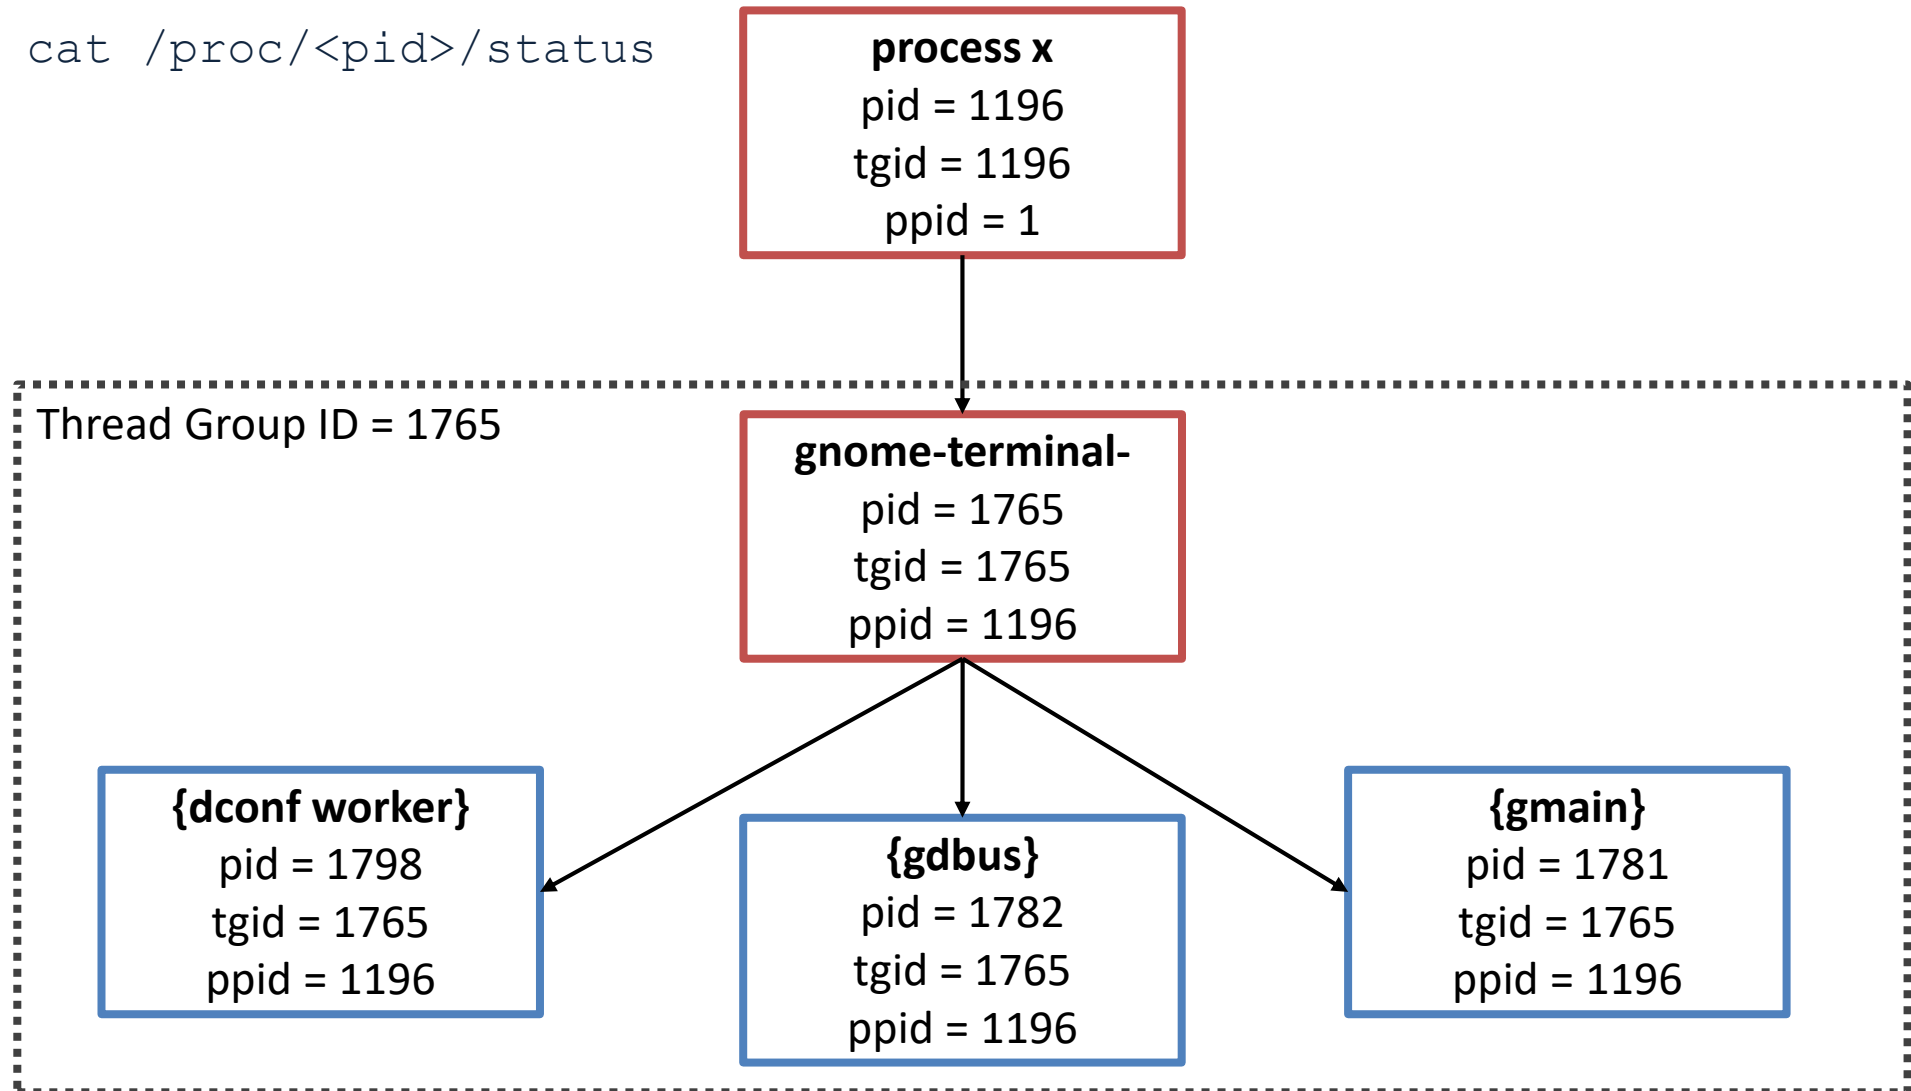
GNOME inter-process communication

GNOME main event loop

- { <name> } → child thread of the parent

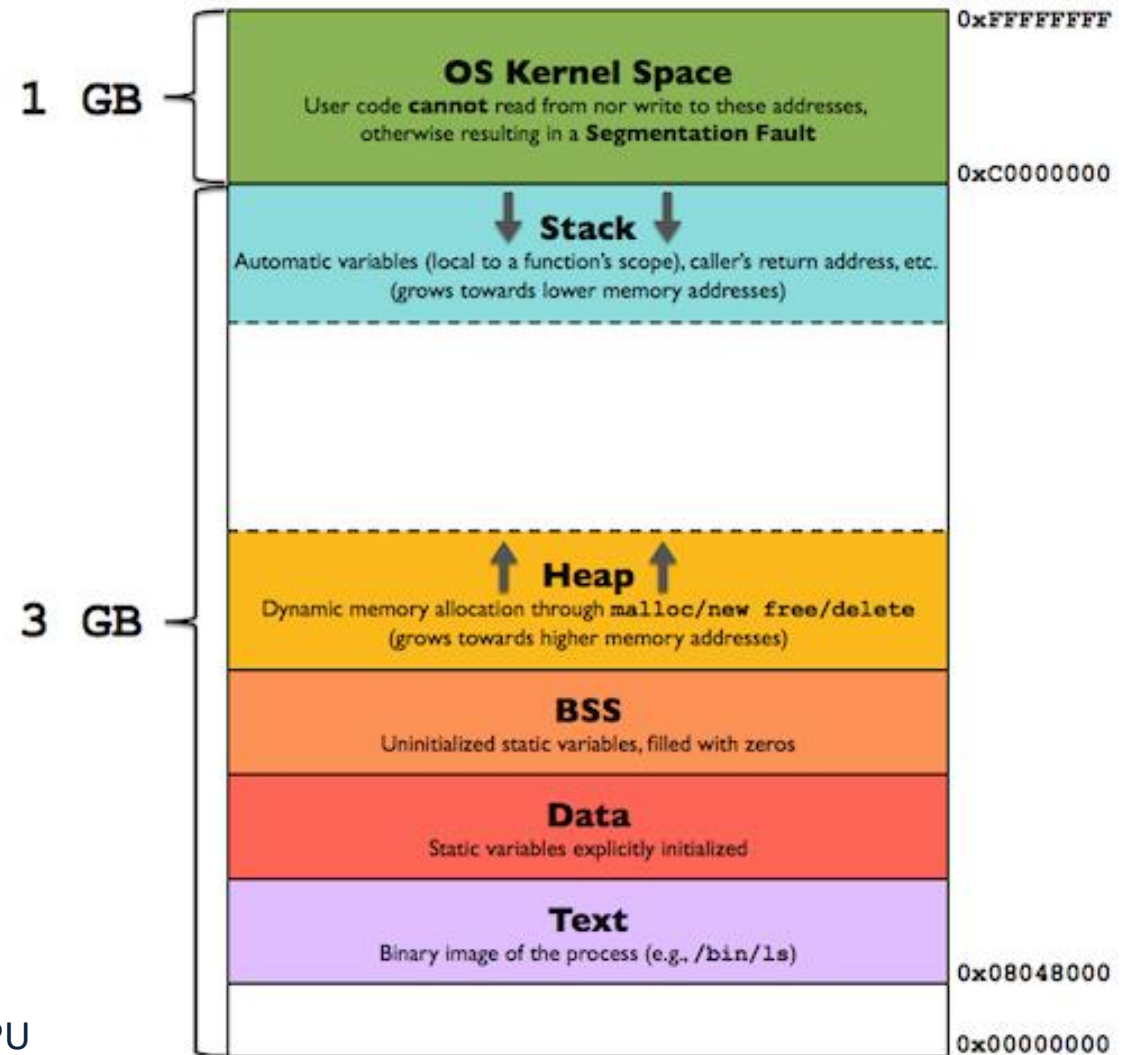
Process IDs

- `cat /proc/<pid>/status`



Program Structures

- Process → a program (= object code) that is executing
 - Program code (.text section)
 - Init. Global variables (.data section)
 - Uninit. Global variables (.bss section)
- Stack
 - Automatic variables (local)
 - Automatically allocated and freed
 - Managed by OS
- Heap
 - Dynamic allocations (global)
 - Managed by program
- OS Kernel Space
 - Shared between processes
 - Used as entrypoints to the kernel for the CPU



Intermezzo: struct

- Composite data type in the C programming language


- Linux Kernel = 97% written in C

```
struct operating_system {  
    int a = 1;  
    char b[10];  
    float *c;  
};
```

- Contiguous block of physical memory
 - Only need one pointer to the start to access all the members
 - Int → 2 or 4 bytes, char → byte array (10 bytes), pointer → 4 or 8 bytes
- Great for data structures, many structs in the kernel

Process Descriptor (1)

- Kernel stores task (=process) list → list of processes → `tasks` vector
 - Circular doubly linked list → item: process descriptor



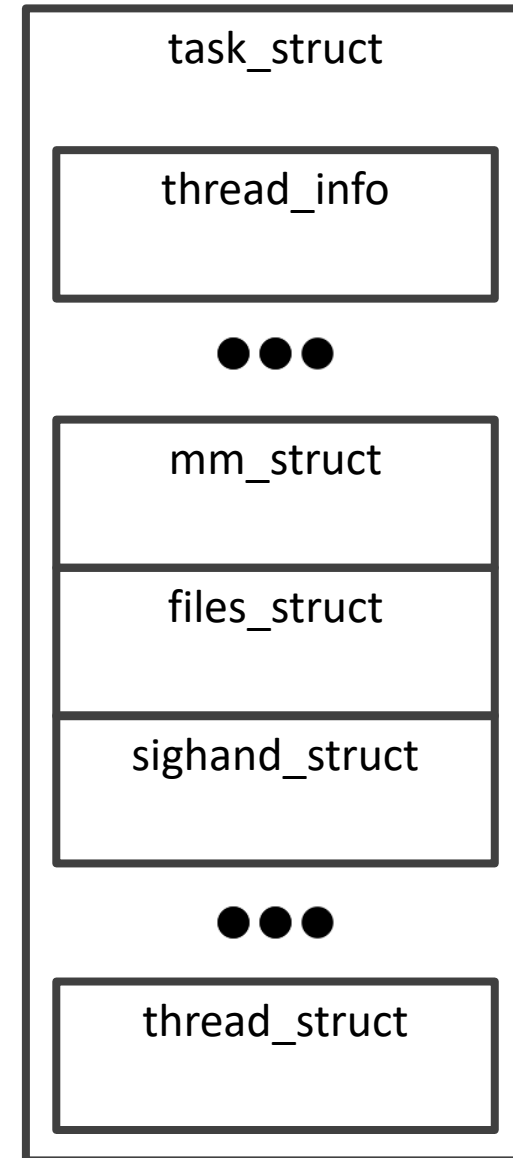
- Process Descriptor → contains all data that describes the executing program
 - `struct task_struct` → pid, priority, state, parent process (`task_struct`)
 - Kernel data structure (exists in kernel space)
- `task_struct` process descriptor
 - <https://github.com/torvalds/linux/blob/master/include/linux/sched.h#L593>

Process Tree

- /sbin/init process → PID = 1
 - Reads init scripts, executes startup programs, completes boot process
- Process has 1 parent
 - Defined in the `task_struct` as `parent` → pointer to another `task_struct`
 - `struct task_struct *real_parent;`
 - Also 0-n children in `task_struct` → pointer to list of `task_struct`
 - `struct list_head children;`
- Ways to iterate over processes:
 - Traverse the tree using parent and children using `real_parent` and `children`
 - Iterate the process list (circular doubly linked list) using `for_each_process()`
 - <https://github.com/torvalds/linux/blob/master/include/linux/sched/signal.h#L565>

task_struct

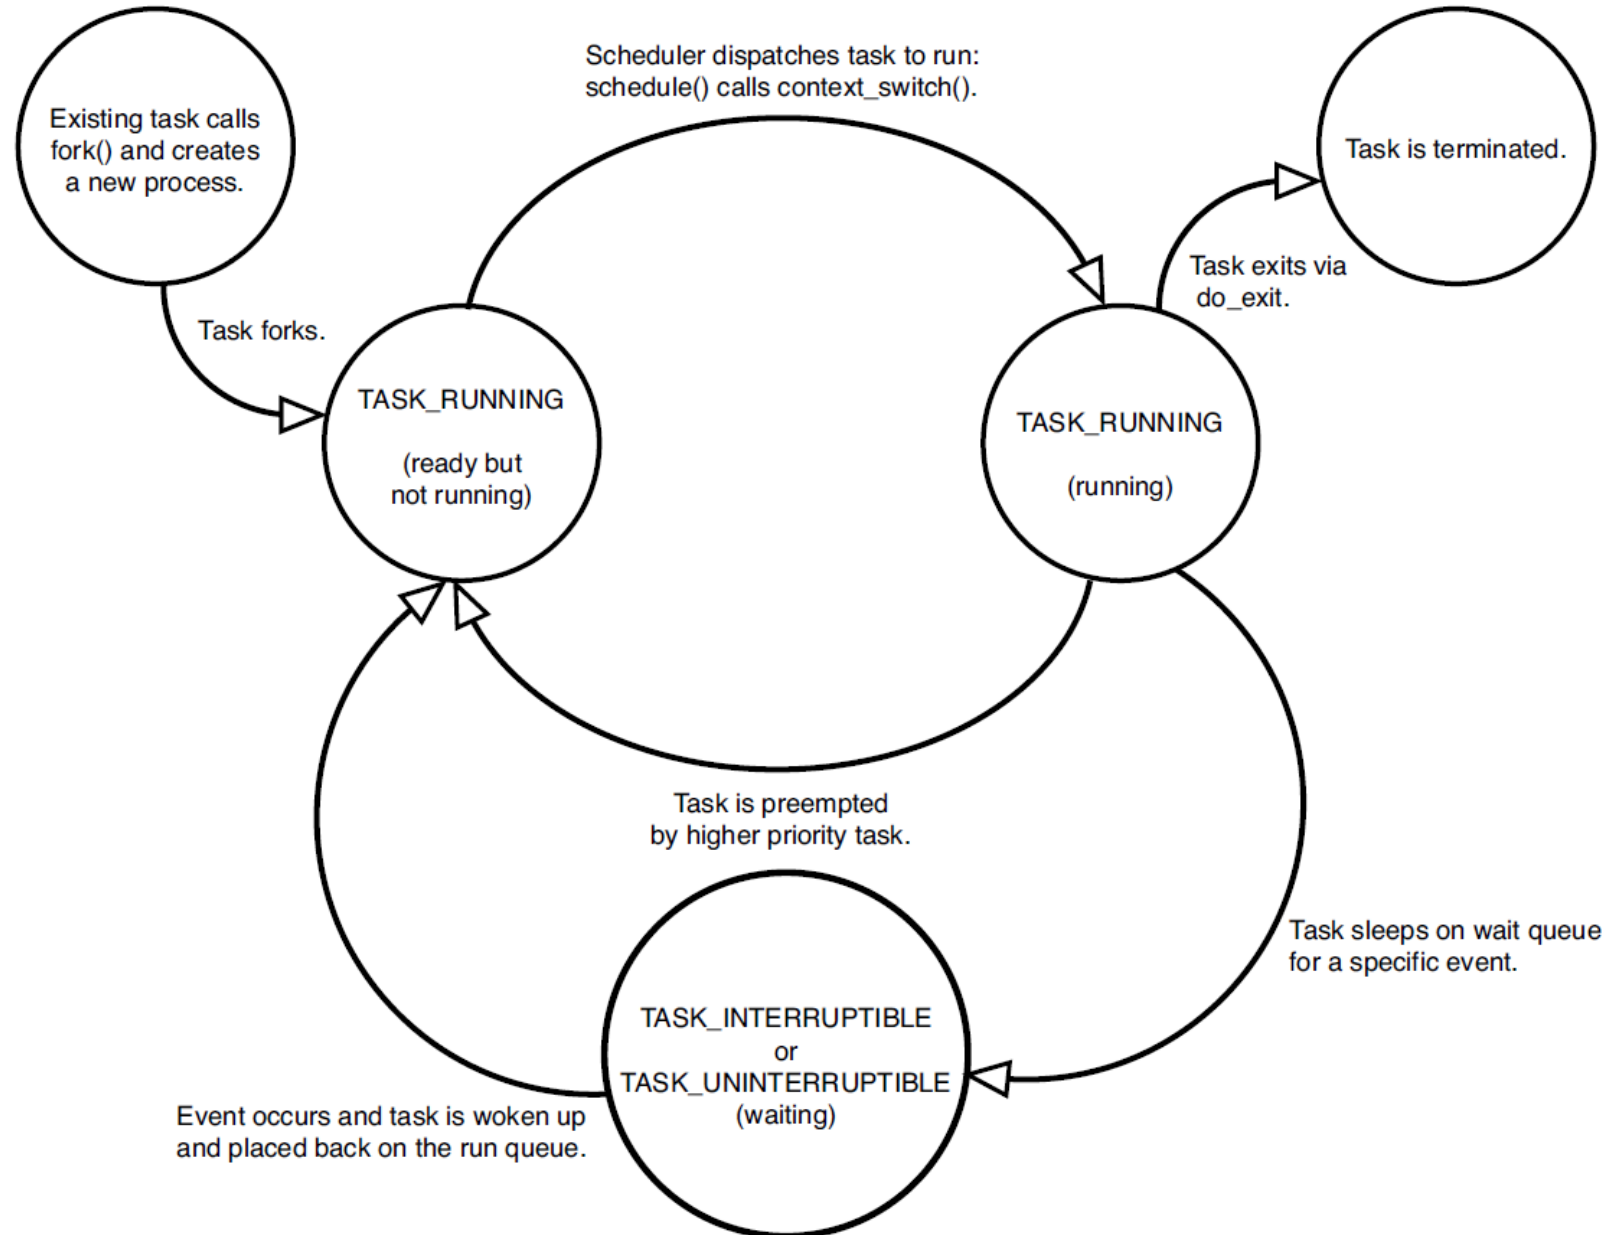
- struct thread_info
 - Low level thread data, pointer to task_struct
- struct mm_struct
 - Memory mappings
- struct files_struct
 - Open files
- struct sighand_struct
 - Signal handlers
- Many, many others
 - 1000's LoC
- struct thread_struct
 - Registers, e.g. stack pointer, data segment



Process Descriptor (2)

- `struct thread_struct`
 - Processor architecture dependent → size of struct depends on processor architecture
 - Therefore this struct is at the end of the `task_struct`
 - <https://github.com/torvalds/linux/tree/master/arch/<processor>/include/asm/processor.h>
 - Some architecture have vastly different features (e.g. more registers)
 - i386 → 16 registers
 - x86-64 → 40 registers

Process State



Process State Flags

- TASK_RUNNING
 - Currently running or waiting in the run-queue
 - All processes in user-space have this state
- TASK_INTERRUPTIBLE
 - Asleep, waiting for some condition, event, or signal
- TASK_UNINTERRUPTIBLE
 - Asleep, but cannot be woken up by signals
- TASK_STOPPED
 - Execution stopped, after receiving signal (e.g. SIGSTOP)
- TASK_ZOMBIE
 - Execution stopped, waiting for parent process to collect exit code

Process Context

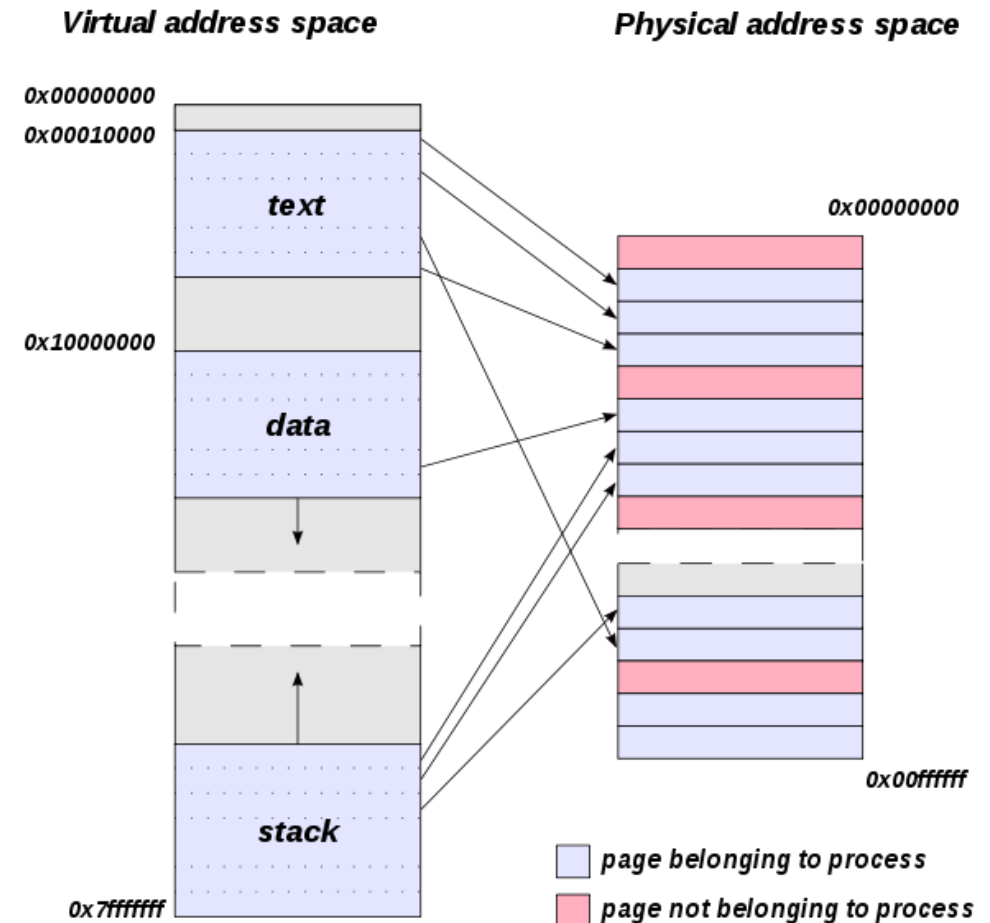
- Process executes program code from executable file
 - Executed within program's address space
 - Occurs in user-space
- Program executes system call or triggers exception
 - Enters kernel-space
 - Kernel executes on behalf of process → “in process context”
- All access to the kernel is through system calls and exception handlers

Process Creation

- Traditional *Spawn* mechanism to create process (`posix_spawn()`)
 - Create new process (descriptor)
 - Create new address space
 - Read executable
 - Begin execution
- Unix
 - Separated in `fork()` and `exec()` functions
 - `fork()` → Create child process that is a full copy (duplicates process) of parent process (PID, PPID, some resources, not copied), does not share resources
 - `exec()` → Loads new executable in address space (replaces process) and begins execution

Process Creation Linux

- Linux Way → Copy-on-write
 - Instead of copying all resources on `fork()`, the copying of data is delayed
- Page → block of contiguous virtual memory with fixed length described by a single entry in the page table
- Page table → per process table that maps virtual to physical addresses
- Page table structure is copied, not pages themselves
- Only when child writes to page, the page is copied (copy-on-write)
- CPU vs Memory trade-off



Child process creation

- To create a child process
- Call `clone()` → `do_fork()` → `dup_task_struct()`
 - Creates new kernel stack, `thread_info` and `task_struct` (identical to parent)
 - Some values in `task_struct` are reset (e.g. stats)
 - Child process state = `TASK_UNINTERRUPTIBLE`
- Call `copy_process()` → `copy_flags()`
 - Update flags for process (e.g. super user privileges flag)
- Call `alloc_pid()`
 - Assign PID
- Share open files, file system information, signal handlers, process address space, namespace (usually shared between threads in same process)
- Return pointer to new child process → `exec()` → child runs first

Threads

- Multiple threads of execution → concurrency and parallelism
 - Share memory address space, files, other resources
- Linux has no explicit concept of threads
 - Thread is just a process that happens to share resources with other processes
 - MS Windows has explicit constructs for threads (lightweight processes)
 - TGID → Thread Group ID (parent PID) and TID → Thread ID
 - `getpid()` always returns the TGID, despite its name!
- Threads created via `clone()`, resource sharing indicated by passed flags
- Kernel Threads → standard processes executing only in kernel space
 - Do not have own address space (shared with kthreads), do not switch to user-space
 - Created only by other kernel threads (parent is `kthreadd`)

Flag	Meaning
<code>CLONE_FILES</code>	Parent and child share open files.
<code>CLONE_FS</code>	Parent and child share filesystem information.
<code>CLONE_IDLETASK</code>	Set PID to zero (used only by the idle tasks).
<code>CLONE_NEWNS</code>	Create a new namespace for the child.
<code>CLONE_PARENT</code>	Child is to have same parent as its parent.
<code>CLONE_PTRACE</code>	Continue tracing child.
<code>CLONE_SETTID</code>	Write the TID back to user-space.
<code>CLONE_SETTLS</code>	Create a new TLS for the child.
<code>CLONE_SIGHAND</code>	Parent and child share signal handlers and blocked signals.
<code>CLONE_SYSVSEM</code>	Parent and child share System V <code>SEM_UNDO</code> semantics.
<code>CLONE_THREAD</code>	Parent and child are in the same thread group.
<code>CLONE_VFORK</code>	<code>vfork()</code> was used and the parent will sleep until the child wakes it.
<code>CLONE_UNTRACED</code>	Do not let the tracing process force <code>CLONE_PTRACE</code> on the child.
<code>CLONE_STOP</code>	Start process in the <code>TASK_STOPPED</code> state.
<code>CLONE_SETTLS</code>	Create a new TLS (thread-local storage) for the child.
<code>CLONE_CHILD_CLEARTID</code>	Clear the TID in the child.
<code>CLONE_CHILD_SETTID</code>	Set the TID in the child.
<code>CLONE_PARENT_SETTID</code>	Set the TID in the parent.
<code>CLONE_VM</code>	Parent and child share address space.

fork() vs. clone()

- `fork()` → POSIX standard
 - Creates a new child process, no shared resources
 - In Linux, `fork()` calls `clone()` with specific parameters
 - `clone(..., SIGCHLD, 0);`
 - Child process execution continues from where `fork()` was called
 - No parameters
- `clone(fn, stack, flags, args)` → Linux specific = not portable
 - Creates a new child process, shared resources (depending on flags), e.g. new thread:
 - `clone(..., CLONE_VM | CLONE_FS | CLONE_FILES | CLONE_SIGHAND | CLONE_THREAD, 0)`
 - Child process execution starts from the function that is passed to `clone()`
 - Many parameters, **note**: stack grows down, set stack pointer to highest address

Process Termination

- Process terminates → kernel releases resources and notifies parent
- Steps
 - In `task_struct` set `PF_EXITING` flag
 - Call `del_timer_sync()` to remove kernel timers
 - Call `exit_mm()` to release `mm_struct` (if not shared → destroyed)
 - Call `exit_sem()` to stop Inter Process Communication (IPC)
 - Call `exit_files()` / `exit_fs()` to decrease file descriptor usage count
 - In `task_struct` set `exit_code`
 - Call `exit_notify()` to notify parent, in `task_struct` set `exit_state` `EXIT_ZOMBIE`
 - Now process is unrunnable (`schedule()` called to reschedule another process), only memory usage is from `task_struct`, `thread_struct`.
 - Wait for parent to collect exit information → remaining memory is freed

Parentless Child

- What if parent process exits before child process? → reparent
 - Another process in same thread group
 - init (PID = 1) process
- This process has to be executed for each (soon-to-be) parentless child
- Example `ptrace()` → process tracing
 - Temporarily changes process parent to the process that is ptracing
 - Parent can exit before child


Useful Terminal Commands

- `ps` command
 - Information about currently running processes
- `top` command
 - Interactive process information
- `ps tree` command
 - Display the process tree for a given process ID
- `pgrep` command
 - Get PID by name
- `kill`, `pkill`, `killall` commands
 - Send (termination) signals to process by PID or name
- `cat /proc/<pid>/status`
 - Print process information

Mini Crash Course Kernel Modules

Introduction to C

- C is a general-purpose, imperative programming language
- 1972 – Dennis Ritchie (also co-creator of Unix) creates C language at Bell Labs
- 1978 – Publication of The C Programming Language book
- 1983 – American Nation Standards Institute begins standardization
- 1988 – ANSI C standard completed
- Key features
 - Structured programming
 - Lexical (static) variable scope
 - Recursion
 - Static type system

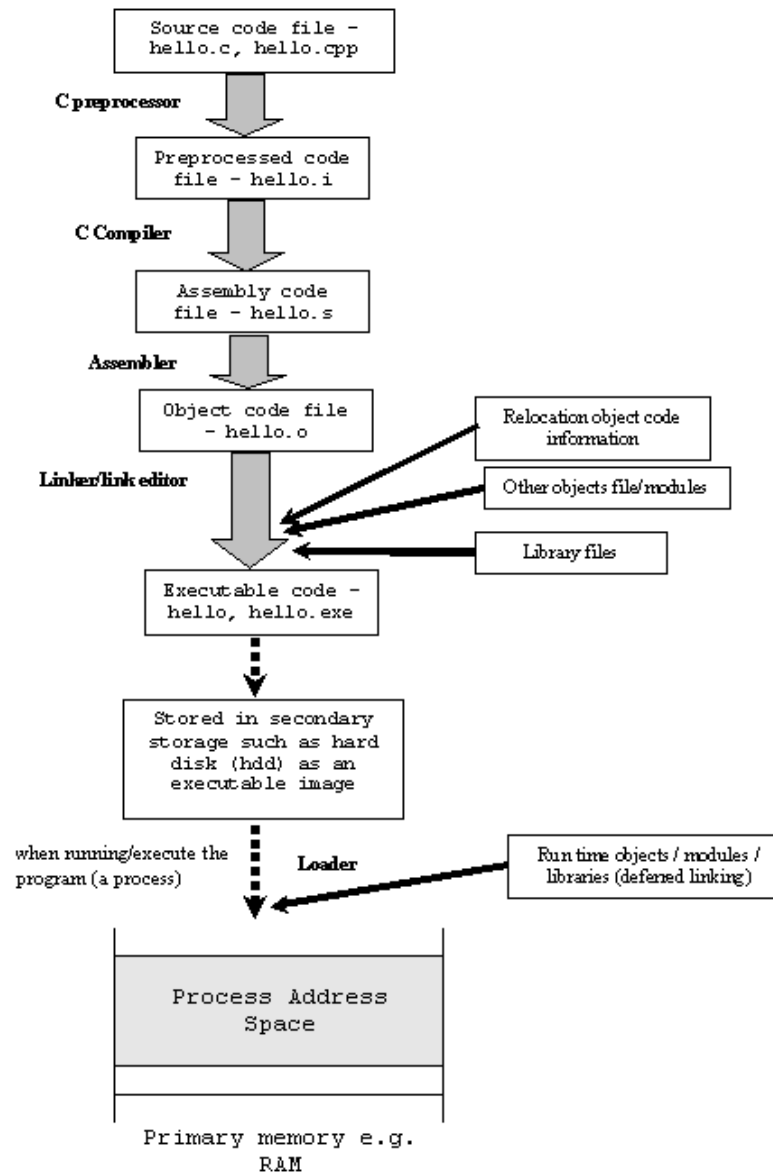


Dennis Ritchie

Pre-processor, compiler, assembler, linker, loader

- GNU C Compiler → `gcc` command
 - `apt install build-essential` (as root)
 - Runs the pre-processor, compiler, runs the assembler, optionally runs the linker
- Pre-processor → include files, conditional compilation, macros
- Compiler → take pre-processed source code, generate assembler code
- Assembler → translate assembler to machine code, generate object file
- Linker → takes object files and libraries, combines into one executable
- Loader → loads the program in memory (`exec()` system call), create process

Pre-processor, compiler, assembler, linker, loader



Hello World

- `hello.c`

```
#include <stdio.h>
```

System header file, declares interface to the system

```
int main( void )
```

Function definition. Main entry point, one per application

```
{
```

```
    printf("Hello World!\n");
```

Write to stdout, declared in `stdio.h`

```
    return 0;
```

Return value of the function, exit code

```
}
```

- To compile: `gcc -w -wError -o hello hello.c`

Compiler flag, all warnings enabled

Compiler flag, warnings to errors

Output executable filename

Files to compile

Makefile

- makefile → organize code compilation, make command
 - Defines a set of tasks to be executed

```
CC=gcc
```

```
CFLAGS=-w -Werror
```

```
TARGET=hello
```

CC, CFLAGS, and TARGET are constants

```
all:
```

Make target named 'all'

```
$(CC) $(CFLAGS) -o $(TARGET) $(TARGET).c
```

Translates to: gcc -w -Werror -o hello hello.c

- Execute `make <target>`, to execute a specific target
- Execute `make`, first target will be executed

Kernel Module

- A kernel module is an object file (.ko) containing code to extend the running kernel
 - Support hardware
 - Add new filesystems
 - Introduce system calls
- Implemented in C
 - Access to kernel space and kernel functions and macros
 - `task_struct`, `for_each_process()`, `thread_info`, etc.
 - No user space libraries (e.g. standard C library)
- Compiled, assembled, but not linked until runtime (by the kernel)

Hello Kernel

■ hello-module.c

```
#include <linux/module.h>
#define M_AUTHOR "Brian Setz <brian.setz@iaas.uni-stuttgart.de>"
#define M_DESC "Hello Module"
int init(void)
{
    printk(KERN_INFO "init(void)\n");
    return 0;
}
void exit(void)
{
    printk(KERN_INFO "exit(void)\n");
}
module_init(init);
module_exit(exit);
MODULE_LICENSE("GPL");
MODULE_AUTHOR(M_AUTHOR);
MODULE_DESCRIPTION(M_DESC);
```

Header file required for all modules

Definitions for module information

printk is kernel's version of printf

If not 0, installation is considered failed

Kernel levels defined in linux/kern_levels.h

Define module entry point

Define module exit point (clean up, etc.)

If not GPL, kernel is marked as tainted

Module information

Compile and Install

- `apt install module-assistant`
 - `m-a prepare` → installs kernel headers, when you update kernel, rerun
- `make` → builds the kernel object file (assuming the Makefile is there)
- `modinfo <x>.ko` → inspects the kernel module
- `insmod <x>.ko` → installs the kernel module
- `cat /var/log/kern.log` → printk logs to the kern.log
- `rmmod <x>` → uninstalls the kernel module


Resources

- Books

- The C Programming Language (ANSI C), Kernighan & Ritchie



- Programming in ANSI C, Balagurusamy (8th edition, 2018)



- Many online resources and courses

- C keywords: <https://www.programiz.com/c-programming/list-all-keywords-c-language>
 - Linux Kernel Module programming: <https://www.tldp.org/LDP/lkmpg/2.6/lkmpg.pdf>

Summary

- Process vs. Thread
- Process Descriptor
- Process State
- Process Creation
- Process Termination
- Mini Crash Course C / Kernel Modules

Exercises

Exercises

Always include a full description (e.g. input / output, steps you take) of the performed tasks, implementation of code, in addition to the answers.

1. Use only the `ps` command to **(a)** list all processes (including UID) with parent process ID (PPID) 1. **(b)** Filter the list of processes such that only processes running as username root remain. **(c)** For the remaining processes, provide a 2-3 sentences describing their functionality.

Download `OS_Lab2.zip` and unzip it.

2. Open the `fork` directory and inspect the files. Compile the source file into a binary using the `make` command. **(a)** Run the binary and show the process tree using `pstree` and the parent PID. **(b)** Explain the number of running processes based on the source code in `fork.c`. **(c)** Do these processes share memory or other resources? Why (not)?
3. Open the `clone` directory. **(a)** Implement an application in C that uses a `fork()`, a `clone()` system call to create a process, and a `clone()` system call to create a thread. **(b)** For each of the create processes / threads print and show the TGID and PPID. **(c)** Explain why the TGIDs and PPIIDs are as shown.
4. Open the `process-module` directory and inspect the files. **(a)** Implement the `list_processes` function such that it outputs a list of all processes including the executable name, process ID, and thread group ID. **(b)** Show the output of the `kern.log` related to your module. **(c)** Which parts of the kernel module execute in kernel space?